

Chapter 2: Process Management

2.1 Processes



GV: Nguyễn Thị Thanh Vân

Operating System Concepts – 10th Edition

Silberschatz, Galvin and Gagne ©2018



Objectives

- Identify the separate components of a process and illustrate how they are represented and scheduled in an operating system.
- Describe how processes are created and terminated in an operating system, including developing programs using the appropriate system calls that perform these operations.
- Describe and contrast interprocess communication using shared memory and message passing.
- Design programs that use pipes and POSIX shared memory to perform interprocess communication.
- Describe client-server communication using sockets and remote procedure calls.
- Design kernel modules that interact with the Linux operating system.



Operating System Concepts – 10th Edition

3.2

Silberschatz, Galvin and Gagne ©2018



Outlines

- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication
- IPC in Shared-Memory Systems
- IPC in Message-Passing Systems
- Examples of IPC Systems
- Communication in Client-Server Systems



Process Concept

- An OS executes a variety of programs that run as a process.
- **Process**
 - a program in execution - must progress in sequential fashion.
 - No parallel execution of instructions of a single process
- One program can be several processes
 - Consider multiple users executing the same program
 - 4 Compiler, Text editor
- The memory layout of a process is typically divided into multiple parts
 - The program code, also called **text section**
 - Current activity including **program counter**, processor registers
 - **Stack** containing temporary data
 - 4 Function parameters,
 - 4 return addresses,
 - 4 local variables
 - **Data section** containing global variables
 - **Heap** containing memory dynamically allocated during run time





Process in Memory

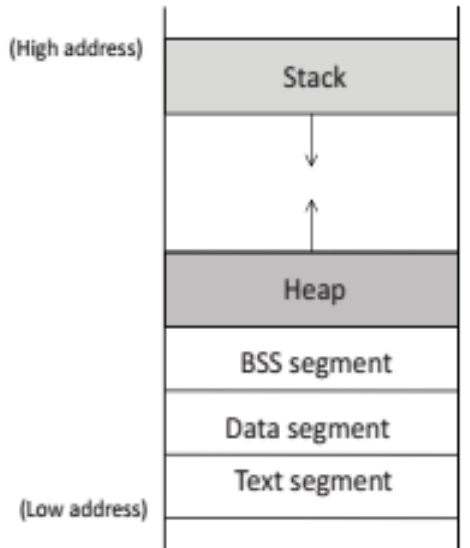
Stack: store local variables in func, store data related to function calls: return address, arguments, (LIFO)

Heap: provide space for dynamic memory allocation. This area is managed by malloc, calloc ...


BSS segment: stores uninitialized static/global variables (zero)

Data segment: stores static/global variables that are initialized by the programmer

Text: stores the executable code of the program (read-only)



Memory Layout of a C Program



```
#include <stdio.h>
#include <stdlib.h>
```

```
int x;
int y = 15;
```

```
int main(int argc, char *argv[])
{
    int *values;
    int i;

    values = (int *)malloc(sizeof(int)*5);

    for(i = 0; i < 5; i++)
        values[i] = i;

    return 0;
}
```






Process

- A process includes:
 - Text
 - Data
 - Heap
 - Stack
 - PC – Program counter: a register that manages the memory address of the instruction to be executed next
 - PSW – Program status word: a register that performs the function of a status register and program counter
 - SP – Stack pointer
 - Registers
- Four principal events cause processes to be created:
 - System initialization.
 - Execution of a process-creation system call by a running process.
 - A user request to create a new process.
 - Initiation of a batch job.



Process State

- As a process executes, it changes **state**
 - **New:** The process is being created
 - **Running:** Instructions are being executed
 - **Waiting:** The process is waiting for some event to occur
 - **Ready:** The process is waiting to be assigned to a processor
 - **Terminated:** The process has finished execution






Process Control Block (PCB)

- ❖ PCB: a data structure used by computer OS to store all the information about a process.
- ❖ Each process is represented in OS by PCB, also called **task control block**
- ❖ It contains many pieces of information associated with a specific process:

- Process state – running, waiting, etc.
- Program counter – location of instruction to next execute
- CPU registers – contents of all process-centric registers
- CPU scheduling information- priorities, scheduling queue pointers
- Memory-management information – memory allocated to the process
- Accounting information – CPU used, clock time elapsed since start, time limits
- I/O status information – I/O devices allocated to process, list of open files



Threads

- So far, process has a single thread of execution
- Most modern OSs have extended the process concept to allow a process to have multiple threads of execution
 - thus to perform more than one task at a time
 - Multiple threads can run in parallel
- The PCB is expanded to include information for each thread.
 - Must then have storage for thread details,
- Explore in detail later – next chapter






PCB in Linux OS

PCB is represented by the C structure `task_struct`, which is found in `<include/linux/sched.h>`

```
pid_t pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice; /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```



Operating System Concepts – 10th Edition

3.12

Silberschatz, Galvin and Gagne ©2018



Process Scheduling

- Goals of:
 - Multiprogramming: some process running at all times so as to maximize CPU utilization
 - Time sharing: switch a CPU core among processes so frequently that users can interact with each program while it is running
- => **Process scheduler** selects among available processes for next execution on CPU core
- The number of processes currently in memory is known as the **degree of multiprogramming**
- Maintains **scheduling queues** of processes
 - **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
 - **Wait queues** – set of processes waiting for an event (i.e., I/O)
 - Processes migrate among the various queues

Operating System Concepts – 10th Edition

3.13


Silberschatz, Galvin and Gagne ©2018





Ready and Wait Queues

- Ready and Wait Queues are generally stored as linked lists

Operating System Concepts – 10th Edition

3.14

Silberschatz, Galvin and Gagne ©2018



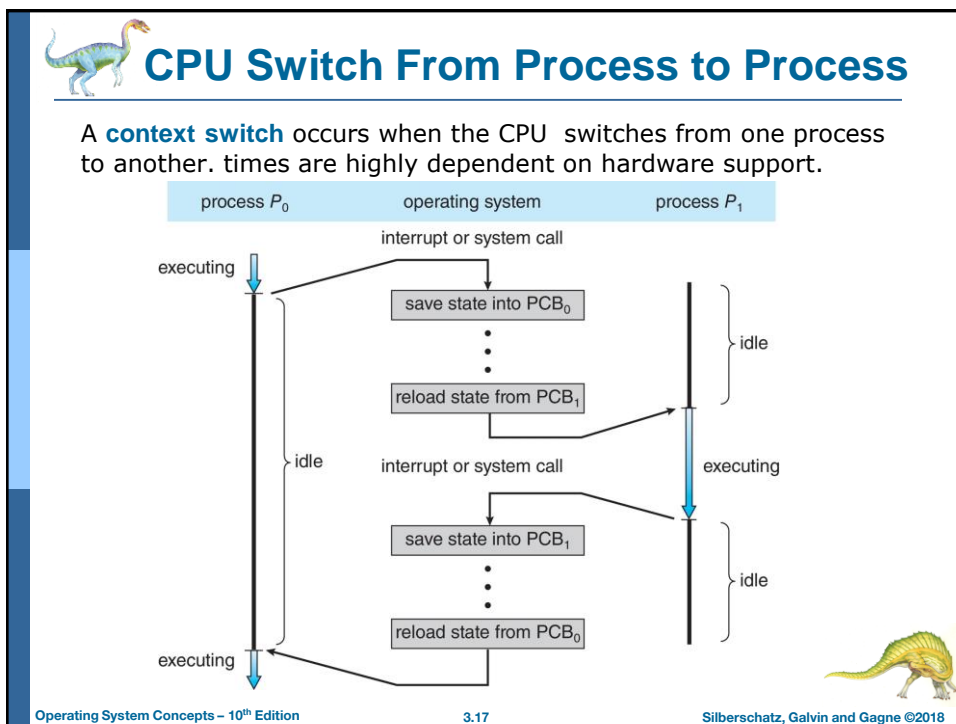
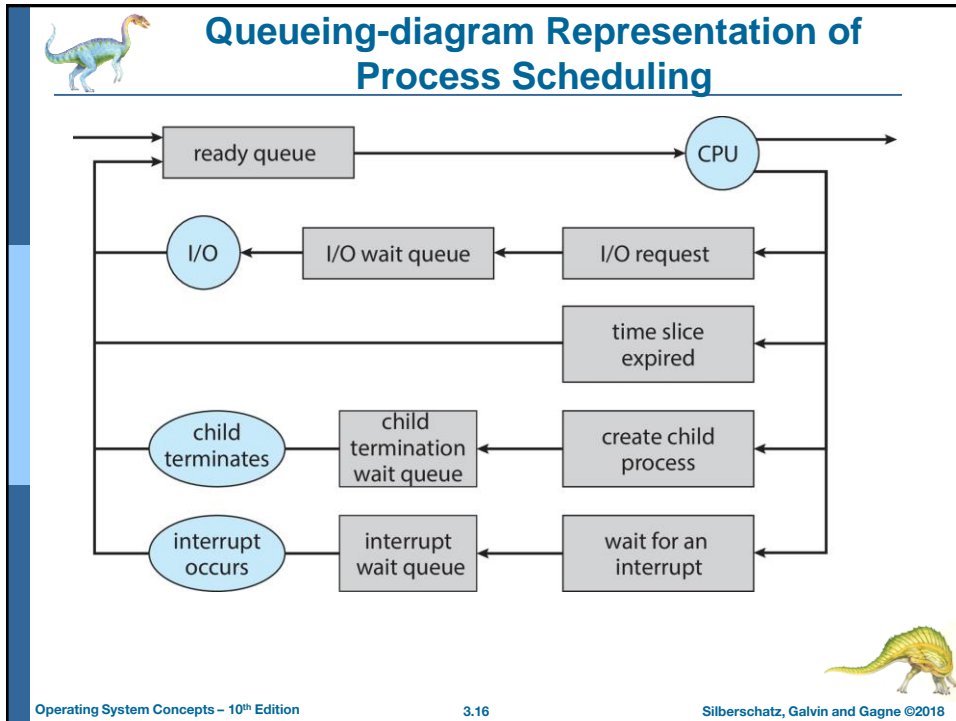
CPU Scheduling

- A process migrates among the ready queue and various wait queues throughout its lifetime.
- CPU scheduler:
 - select from among the processes that are in the ready queue and allocate a CPU core to one of them
 - select a new process for the CPU frequently.
- Queueing-diagram Representation of Process Scheduling: Once the process is allocated the CPU and is executing, one of several events could occur:
 - The process could issue an I/O request and then be placed in an I/O queue.
 - The process could create a new child process and wait for the child's termination.
 - The process could be removed forcibly from the CPU, as a result of an interrupt, and be put back in the ready queue.

Operating System Concepts – 10th Edition

3.15

Silberschatz, Galvin and Gagne ©2018





Context Switching

- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process represented in the PCB
- Context-switch time is pure overhead; the system does no useful work while switching
 - The more complex the OS and the PCB, the longer the context switch
- Time dependent on hardware support
 - Some hardware provides multiple sets of registers per CPU, multiple contexts loaded at once



Multitasking in Mobile Systems

- Some mobile systems (e.g., early version of iOS) allow only one process to run, others suspended
- Due to screen real estate, user interface limits iOS provides for a
 - Single **foreground** process- controlled via user interface
 - Multiple **background** processes- in memory, running, but not on the display, and with limits
 - Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback
- Android runs foreground and background, with fewer limits
 - Background process uses a **service** to perform tasks
 - Service can keep running even if background process is suspended
 - Service has no user interface, small memory use





Operations on Processes

- System must provide mechanisms for:
 - Process creation
 - Process termination



Process Creation


- **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- Generally, process identified and managed via a **process identifier (pid)**
- Resource sharing options
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate





Process Creation (Cont.)

- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - **fork()** system call creates new process
 - **exec()** system call used after a **fork()** to replace the process' memory space with a new program
 - Parent process calls **wait()** waiting for the child to terminate


Operating System Concepts – 10th Edition

3.22

Silberschatz, Galvin and Gagne ©2018



A Tree of Processes in Linux

Operating System Concepts – 10th Edition

3.23

Silberschatz, Galvin and Gagne ©2018





C Program Forking Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```



Creating a Separate Process via Windows API

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```





Process Termination

- Process executes last statement and then asks the operating system to delete it using the **exit()** system call.
 - Returns status data from child to parent (via **wait()**)
 - Process' resources are deallocated by operating system
- Parent may terminate the execution of children processes using the **abort()** system call. Some reasons for doing so:
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting, and the operating systems does not allow a child to continue if its parent terminates



Process Termination

- Some operating systems do not allow child to exists if its parent has terminated. If a process terminates, then all its children must also be terminated.
 - **cascading termination.** All children, grandchildren, etc., are terminated.
 - The termination is initiated by the operating system.
- The parent process may wait for termination of a child process by using the **wait()** system call. The call returns status information and the pid of the terminated process

pid = wait(&status);

- If no parent waiting (did not invoke **wait()**) process is a **zombie**
- If parent terminated without invoking **wait()**, process is an **orphan**





Ex

- Detail in Process_syscal_exp1 file



Android Process Importance Hierarchy


- Mobile operating systems often have to terminate processes to reclaim system resources such as memory. From **most** to **least** important:
 - Foreground process
 - Visible process
 - Service process
 - Background process
 - Empty process
- Android will begin terminating processes that are least important.





Multiprocess Architecture – Chrome Browser

- Many web browsers ran as single process (some still do)
 - If one web site causes trouble, entire browser can hang or crash
- Google Chrome Browser is multiprocess with 3 different types of processes:
 - **Browser** process manages user interface, disk and network I/O
 - **Renderer** process renders web pages, deals with HTML, Javascript. A new renderer created for each website opened
 - 4 Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
 - **Plug-in** process for each type of plug-in



Each tab represents a separate process.



Interprocess Communication


- Processes within a system may be **independent** or **cooperating**
- **Cooperating process** can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience
- Cooperating processes need **interprocess communication (IPC)**
- Two models of IPC
 - **Shared memory** (under the control of users)
 - **Message passing** (under the control of OS)






Communications Models

(a) Shared memory.



(a)

(b) Message passing.




(b)



Producer-Consumer Problem

- Paradigm for cooperating processes:
 - producer** process produces information that is consumed by a **consumer** process




- Two variations:
 - unbounded-buffer** places no practical limit on the size of the buffer:
 - 4 Producer never waits
 - 4 Consumer waits if there is no buffer to consume
 - bounded-buffer** assumes that there is a fixed buffer size
 - 4 Producer must wait if all buffers are full
 - 4 Consumer waits if there is no buffer to consume





Shared Memory Solution

- An area of memory shared among the processes that wish to communicate
- The communication is under the control of the users processes not the operating system.
- Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory.
- Synchronization is discussed in great details in Chapters 6 & 7.

Operating System Concepts – 10th Edition

3.34

Silberschatz, Galvin and Gagne ©2018




Bounded-Buffer – Shared-Memory Solution

- Shared data


```

#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
      
```



- Solution presented in next slides is correct, but can only use **BUFFER_SIZE-1** items; that is: 9 items

Operating System Concepts – 10th Edition

3.35

Silberschatz, Galvin and Gagne ©2018



Producer/ Consumer Process – Shared Memory

Producer

```
item next_produced;
while (true) {
    /* produce an item in next produced */
    while ((in + 1) % BUFFER_SIZE == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

Customer

```
item next_consumed;
while (true) {
    while (in == out)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    /* consume the item in next consumed */
}
```



What about Filling all the Buffers?

- Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers.
- We can do so by having an integer **counter** that keeps track of the number of full buffers.
- Initially, **counter** is set to 0.
- The integer **counter** is incremented by the producer after it produces a new buffer.
- The integer **counter** is and is decremented by the consumer after it consumes a buffer.





Producer/ Consumer Process

Producer

```
while (true) {
    /* produce an item in next produced
    while (counter == BUFFER_SIZE)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
    counter++;
}
```

Customer

```
while (true) {
    while (counter == 0)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--;
    /* consume the item in next consumed */
}
```



Race Condition

- **counter++** could be implemented as


```
register1 = counter
register1 = register1 + 1
counter = register1
```
- **counter--** could be implemented as


```
register2 = counter
register2 = register2 - 1
counter = register2
```
- Consider this execution interleaving with “count = 5” initially:


S0: producer execute	register1 = counter	{register1 = 5}
S1: producer execute	register1 = register1 + 1	{register1 = 6}
S2: consumer execute	register2 = counter	{register2 = 5}
S3: consumer execute	register2 = register2 - 1	{register2 = 4}
S4: producer execute	counter = register1	{counter = 6}
S5: consumer execute	counter = register2	{counter = 4}
- **Question** – why was there no race condition in the first solution (where at most N – 1) buffers can be filled?





IPC – Message Passing

- Processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - **send**(message)
 - **receive**(message)
- The *message* size is either fixed or variable



Message Passing (Cont.)

- If processes *P* and *Q* wish to communicate, they need to:
 - Establish a **communication link** between them
 - Exchange messages via send/receive
- Implementation issues:
 - How are **links** established?
 - Can a link be associated with more than two processes?
 - How many links can there be between every pair of communicating processes?
 - What is the capacity of a link?
 - Is the size of a message that the link can accommodate fixed or variable?
 - Is a link unidirectional or bi-directional?





Implementation of Communication Link

- Physical:
 - Shared memory
 - Hardware bus
 - Network
- Logical:
 - Direct or indirect
 - Synchronous or asynchronous
 - Automatic or explicit buffering



Direct Communication


- Processes must name each other explicitly:
 - **send**(P , *message*) – send a message to process P
 - **receive**(Q , *message*) – receive a message from process Q
- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional





Indirect Communication

- Messages are directed and received from mailboxes (also is ports)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional



Indirect Communication

- Operations
 - Create a new mailbox (port)
 - Send and receive messages through mailbox
 - Delete a mailbox
- Primitives are defined as:
 - **Send**(A, message) – send a message to mailbox A
 - **receive**(A, message) – receive a message from mailbox A





Indirect Communication (Cont.)

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A
 - P_1 sends; P_2 and P_3 receive
 - Who gets the message?
- **Solutions**
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a receive operation
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

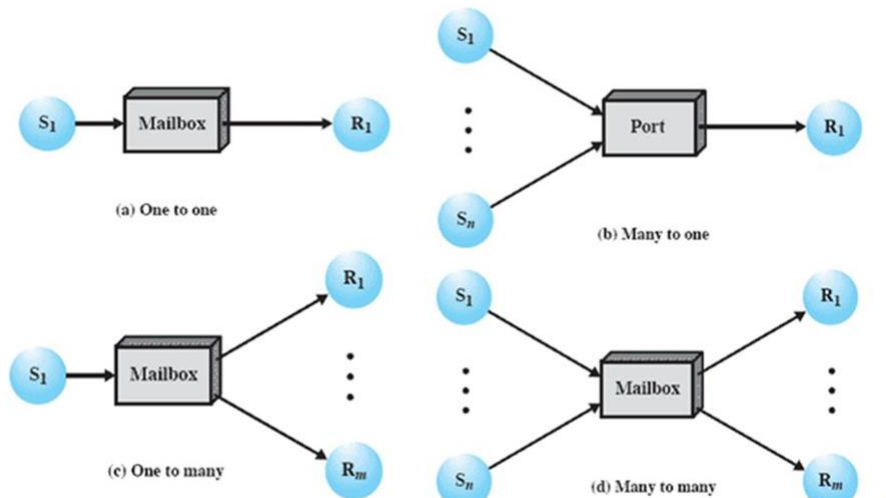
Operating System Concepts – 10th Edition

3.46

Silberschatz, Galvin and Gagne ©2018



Indirect Communication (Cont.)

Operating System Concepts – 10th Edition

3.47

Silberschatz, Galvin and Gagne ©2018



Synchronization

Message passing may be either blocking or non-blocking

- **Blocking** is considered **synchronous**
 - **Blocking send** -- the sender is blocked until the message is received
 - **Blocking receive** -- the receiver is blocked until a message is available
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking send** -- the sender sends the message and continue
 - **Non-blocking receive** -- the receiver receives:
 - 4 A valid message, or
 - 4 Null message
- Different combinations possible
 - If both send and receive are blocking, we have a **rendezvous**



Producer-Consumer: Message Passing

- Producer


```
message next_produced;
while (true) {
    /* produce an item in next_produced */

    send(next_produced);
}
```
- Consumer


```
message next_consumed;
while (true) {
    receive(next_consumed)

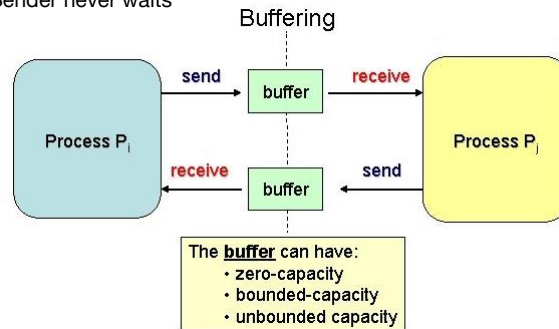
    /* consume the item in next_consumed */
}
```





Buffering

- Queue of messages attached to the link.
- Implemented in one of three ways
 1. Zero capacity – no messages are queued on a link.
Sender must wait for receiver (rendezvous)
 2. Bounded capacity – finite length of n messages
Sender must wait if link full
 3. Unbounded capacity – infinite length
Sender never waits

Operating System Concepts – 10th Edition

3.50

Silberschatz, Galvin and Gagne ©2018



Examples of IPC Systems - POSIX

- POSIX Shared Memory
 - Process first creates shared memory segment
`shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);`
 - Also used to open an existing segment
 - Set the size of the object
`ftruncate(shm_fd, 4096);`
 - Use `mmap()` to memory-map a file pointer to the shared memory object
 - Reading and writing to shared memory is done by using the pointer returned by `mmap()`.

Operating System Concepts – 10th Edition

3.51

Silberschatz, Galvin and Gagne ©2018





IPC POSIX Producer

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr, "%s", message_0);
    ptr += strlen(message_0);
    sprintf(ptr, "%s", message_1);
    ptr += strlen(message_1);

    return 0;
}
```



IPC POSIX Consumer

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```





Examples of IPC Systems - Mach

- Mach communication is message based
 - Even system calls are messages
 - Each task gets two ports at creation- Kernel and Notify
 - Messages are sent and received using the `mach_msg()` function
 - Ports needed for communication, created via `mach_port_allocate()`
 - Send and receive are flexible, for example four options if mailbox full:
 - 4 Wait indefinitely
 - 4 Wait at most n milliseconds
 - 4 Return immediately
 - 4 Temporarily cache a message



Mach Messages

```
#include<mach/mach.h>

struct message {
    mach_msg_header_t header;
    int data;
};

mach_port_t client;
mach_port_t server;
```





Mach Message Passing – Client/Server

```

/* Client Code */

struct message message;

// construct the header
message.header.msgh_size = sizeof(message);
message.header.msgh_remote_port = server;
message.header.msgh_local_port = client;

// send the message
mach_msg(&message.header, // message header
MACH_SEND_MSG, // sending a message
sizeof(message), // size of message sent
0, // maximum size of received message - unnecessary
MACH_PORT_NULL, // name of receive port - unnecessary
MACH_MSG_TIMEOUT_NONE, // no time outs
MACH_PORT_NULL // no notify port
);

/* Server Code */

struct message message;

// receive the message
mach_msg(&message.header, // message header
MACH_RCV_MSG, // sending a message
0, // size of message sent
sizeof(message), // maximum size of received message
server, // name of receive port
MACH_MSG_TIMEOUT_NONE, // no time outs
MACH_PORT_NULL // no notify port
);

```

Operating S

Galvin and Gagne ©2018



Examples of IPC Systems – Windows

- Message-passing centric via **advanced local procedure call (LPC)** facility
 - Only works between processes on the same system
 - Uses ports (like mailboxes) to establish and maintain communication channels
 - Communication works as follows:
 - 4 The client opens a handle to the subsystem's **connection port** object.
 - 4 The client sends a connection request.
 - 4 The server creates two private **communication ports** and returns the handle to one of them to the client.
 - 4 The client and server use the corresponding port handle to send messages or callbacks and to listen for replies.

Operating System Concepts – 10th Edition

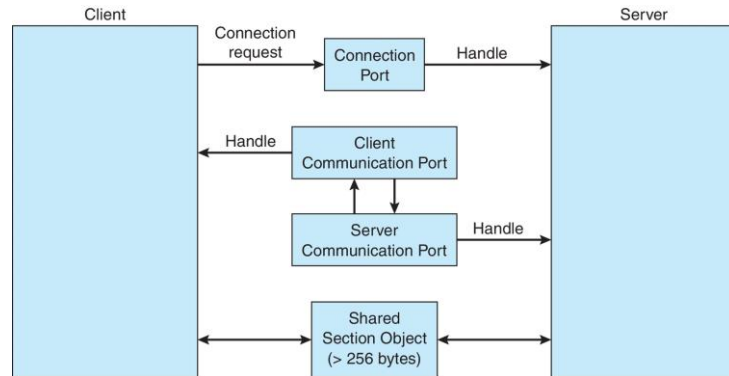
3.57

Silberschatz, Galvin and Gagne ©2018






Local Procedure Calls in Windows



Pipes

- Acts as a conduit allowing two processes to communicate




- Issues:
 - Is communication unidirectional or bidirectional?
 - In the case of two-way communication, is it half or full-duplex?
 - Must there exist a relationship (i.e., **parent-child**) between the communicating processes?
 - Can the pipes be used over a network?
- Ordinary pipes** – cannot be accessed from outside the process that created it. Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.
- Named pipes** – can be accessed without a parent-child relationship.





Ordinary Pipes

- Ordinary Pipes allow communication in standard producer-consumer style
- Producer writes to one end (the **write-end** of the pipe)
- Consumer reads from the other end (the **read-end** of the pipe)
- Ordinary pipes are therefore unidirectional
- Require parent-child relationship between communicating processes




- Windows calls these **anonymous pipes**



Named Pipes


- Named Pipes are more powerful than ordinary pipes
- Communication is bidirectional
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows systems






Communications in Client-Server Systems

- Sockets



- Remote Procedure Calls

Operating System Concepts – 10th Edition


3.62

Silberschatz, Galvin and Gagne ©2018



Sockets

- A **socket** is defined as an endpoint for communication
- Concatenation of IP address and **port**
 - port is a number included at start of message packet to differentiate network services on a host
- The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets
- All ports below 1024 are **well known**, used for standard services
- Special IP address 127.0.0.1 (**loopback**) to refer to system on which process is running
- Communication:

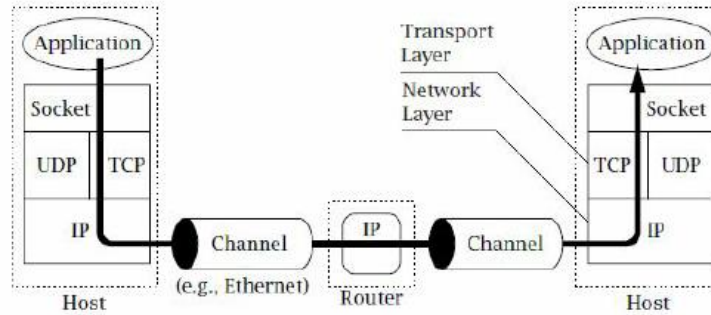
Operating System Concepts – 10th Edition

3.63

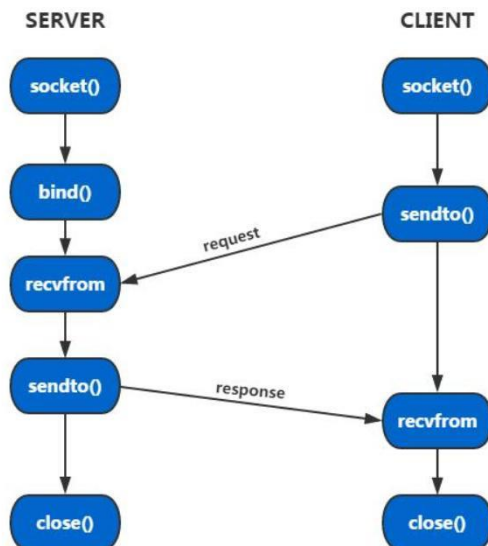
Silberschatz, Galvin and Gagne ©2018



Sockets

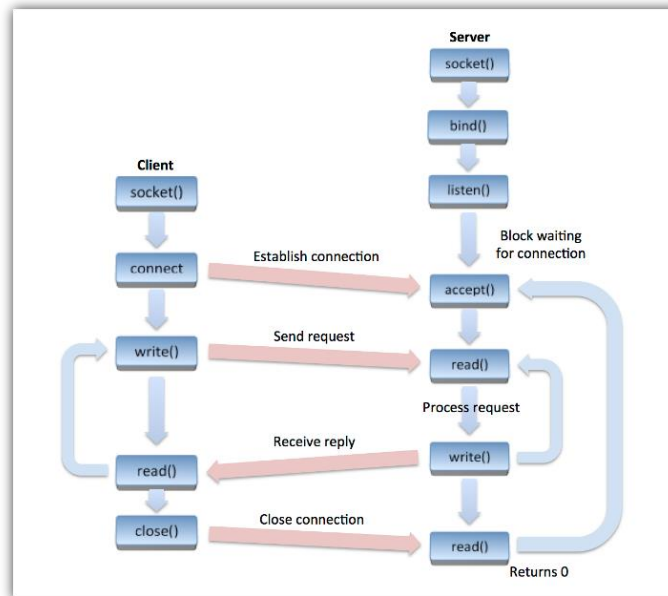


UDP Socket





TCP Socket

Operating System Concepts – 10th Edition

3.66

Silberschatz, Galvin and Gagne ©2018



Sockets in Java

- Three types of sockets
 - **Connection-oriented (TCP)**
 - **Connectionless (UDP)**
 - **MulticastSocket** class—data can be sent to multiple recipients
- Consider this “Date” server in Java:

```

import java.net.*;
import java.io.*;

public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());

                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
  
```

Operating System Concepts – 10th Edition

3.67

Silberschatz, Galvin and Gagne ©2018



Sockets in Java

The equivalent Date client

```
import java.net.*;
import java.io.*;

public class DateClient
{
    public static void main(String[] args) {
        try {
            /* make connection to server socket */
            Socket sock = new Socket("127.0.0.1", 6013);

            InputStream in = sock.getInputStream();
            BufferedReader bin = new
                BufferedReader(new InputStreamReader(in));


            /* read the date from the socket */
            String line;
            while ( (line = bin.readLine()) != null)
                System.out.println(line);

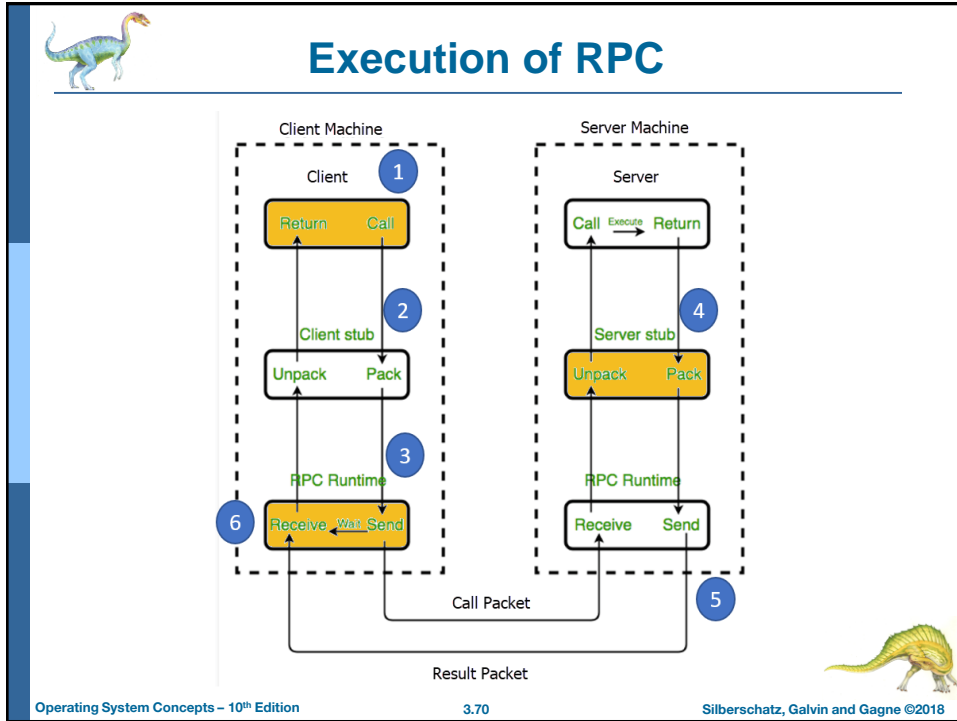
            /* close the socket connection */
            sock.close();
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```



Remote Procedure Calls

- RPC abstracts procedure calls between processes on networked systems
 - Again, uses ports for service differentiation
- Stubs** – client-side proxy for the actual procedure on the server
- The client-side stub locates the server and **marshalls** the parameters
- The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server





End of Chapter 2.1

