

## Chapter 2: Process Synchronization

### 2.5: Deadlocks



GV: Nguyễn Thị Thanh Vân

Operating System Concepts – 10<sup>th</sup> Edition

Silberschatz, Galvin and Gagne ©2018



### Outline

- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock



Operating System Concepts – 10<sup>th</sup> Edition

8.2

Silberschatz, Galvin and Gagne ©2018



## Chapter Objectives

- Illustrate how deadlock can occur when mutex locks are used
- Define the four necessary conditions that characterize deadlock
- Identify a deadlock situation in a resource allocation graph
- Evaluate the four different approaches for preventing deadlocks
- Apply the banker's algorithm for deadlock avoidance
- Apply the deadlock detection algorithm
- Evaluate approaches for recovering from deadlock



## System Model

- System consists of resources
- Resource types  $R_1, R_2, \dots, R_m$ 
  - CPU cycles, memory space, I/O devices
- Each resource type  $R_i$  has  $W_i$  instances.
- Each process utilizes a resource as follows:
  - **Request:** The thread requests the resource
  - **Use:** The thread can operate on the resource
  - **Release:** The thread releases the resource.





## Deadlock Example with Semaphores


- Data:
  - A semaphore `s1` initialized to 1
  - A semaphore `s2` initialized to 1
- Two processes P1 and P2
- P1:

```
wait(s1)
wait(s2)
```
- P2:

```
wait(s2)
wait(s1)
```



## Deadlock Example on a one-way Bridge





## Deadlock Characterization

Deadlock can arise if four conditions hold simultaneously.


- **Mutual exclusion:** only one process at a time can use a resource
- **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes
- **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task
- **Circular wait:** there exists a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2$ , ...,  $P_{n-1}$  is waiting for a resource that is held by  $P_n$ , and  $P_n$  is waiting for a resource that is held by  $P_0$ .



## Resource-Allocation Graph

A set of vertices  $V$  and a set of edges  $E$ .



- $V$  is partitioned into two types:
  - $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system
  - $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system
- **request edge** – directed edge  $P_i \rightarrow R_j$
- **assignment edge** – directed edge  $R_j \rightarrow P_i$





## Resource Allocation Graph Example

- One instance of  $R_1$  ; Two instances of  $R_2$  ; One instance of  $R_3$  ; Three instance of  $R_4$
- $T_1$  holds one instance of  $R_2$  and is waiting for an instance of  $R_1$
- $T_2$  holds one instance of  $R_1$ , one instance of  $R_2$ , and is waiting for an instance of  $R_3$
- $T_3$  holds one instance of  $R_3$



Deadlock

- $T_1 \rightarrow R_1 \rightarrow T_2 \rightarrow R_3 \rightarrow T_3 \rightarrow R_2 \rightarrow T_1$
- $T_2 \rightarrow R_3 \rightarrow T_3 \rightarrow R_2 \rightarrow T_2$

Operating System Concepts – 10<sup>th</sup> Edition

8.9

Silberschatz, Galvin and Gagne ©2018




## Graph with a Cycle But no Deadlock

- $T_1 \rightarrow R_1 \rightarrow T_3 \rightarrow R_2 \rightarrow T_1$
- Observe that:
  - Thread  $T_4$  (holding one instance of  $R_2$ ) may **release its instance of resource type  $R_2$** . That resource can then be allocated to  $T_3$ , breaking the cycle  $\Rightarrow$  no deadlock

### In summary

- If graph contains no cycles  $\Rightarrow$  no deadlock
- If graph contains a cycle  $\Rightarrow$ 
  - If only one instance per resource type, then deadlock
  - If several instances per resource type, possibility of deadlock

$\Rightarrow$  This observation is important when we deal with the deadlock problem.



Operating System Concepts – 10<sup>th</sup> Edition

8.10

Silberschatz, Galvin and Gagne ©2018



## Methods for Handling Deadlocks


- Ensure that the system will **never** enter a deadlock state:
  - Deadlock prevention
    - **Mutual Exclusion:** One or more than one resource are non-shareable (Only one process can use at a time)
    - **Hold and Wait:** A process is holding at least one resource and waiting for resources.
    - **No Preemption:** A resource cannot be taken from a process unless the process releases the resource.
    - **Circular Wait:** A set of processes are waiting for each other in circular form.
  - Deadlock avoidance
- Allow the system to enter a deadlock state and then recover
- Ignore the problem and pretend that deadlocks never occur in the system.
  - What is the consequence?



## Deadlock Prevention

Invalidate one of the four necessary conditions for deadlock:

- **Mutual Exclusion** –
  - sharable resources (e.g., read-only files): not required
  - non-sharable resources (e.g: printer): do not
- **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
  - Require processes to request and be allocated all its resources before it begins execution, or
  - Allow processes to request resources only when the process has none allocated to it.
  - Limit: Low resource utilization; starvation possible





## Deadlock Prevention (Cont.)

### ■ No Preemption:

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all the resources currently being held by the process are released
- Preempted resources are added to the list of resources for which the process is waiting
- Process will be **restarted** only when it can regain its old resources, as well as the new ones that it is requesting

### ■ Circular Wait:

- Impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration




## Circular Wait

- Cách 1: mỗi process yêu cầu thực thể của tài nguyên theo thứ tự tăng dần (định nghĩa bởi hàm  $F$ ) của loại tài nguyên. Ví dụ
  - Chuỗi yêu cầu thực thể hợp lệ: tape drive  $\rightarrow$  disk drive  $\rightarrow$  printer
  - Chuỗi yêu cầu thực thể không hợp lệ: disk drive  $\rightarrow$  tape drive
- Cách 2: Khi một process yêu cầu một thực thể của loại tài nguyên  $R_j$  thì nó phải trả lại các tài nguyên  $R_i$  với  $F(R_i) > F(R_j)$ .
- Ví dụ

“Chứng minh” cho cách 1: phản chứng

- P1:  $F(R_4) < F(R_1)$
- P2:  $F(R_1) < F(R_2)$
- P3:  $F(R_2) < F(R_3)$
- P4:  $F(R_3) < F(R_4)$
- Vậy  $F(R_4) < F(R_4)$ , mâu thuẫn!





## Circular Wait

- Invalidating the circular wait condition is most commonly used.
- Simply assign each resource (i.e., mutex locks) a unique number.
- Resources must be acquired in order.
- Example: two mutex locks

```
first_mutex = 1
second_mutex = 5
```

code for `thread_two` could not be written as follows:

```
/* thread_one runs in this function */
void *do_work_one(void *param)
{
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);

    pthread_exit(0);
}

/* thread_two runs in this function */
void *do_work_two(void *param)
{
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);

    pthread_exit(0);
}
```



## Deadlock Avoidance

Requires that the system has some additional *a priori* information available

- Simplest and most useful model requires that each process declare the **maximum number** of resources of each type that it may need
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition
- Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes







## Safe State

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state
- System is in **safe state** if there exists a sequence  $P_1, P_2, \dots, P_n$  of ALL the processes in the systems such that for each  $P_i$  the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all the  $P_j$  with  $j < i$
- That is:
  - If  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished
  - When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate
  - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on



## Safe State ex

- Ví dụ: Hệ thống có **12** tape drive và 3 quá trình P0, P1, P2

- Tại thời điểm t0, dữ liệu cho như bảng,
  - => hệ thống còn **3** tape drive sẵn sàng (12-5-2-2).
  - Sequence (P1, P0, P2 ): safe:
    - P1: get and return => avail: **3**+2=5
    - P0: get and return => avail: 5+5=10
    - P2: get and return => avail: 10+2=**12**

=> system is in a **safe state**

- Tại thời điểm t1, P2 yêu cầu và được cấp phát 1 tape drive còn **2** tape drive sẵn sàng
  - P1: get and return => avail: **2**+2=4
  - P0: get 5: wait
  - P2: get

=> Sequence (P1, P0, P2 ): no safe

=> system is in a no safe state

Process	Max - need	Allocation
P0	10	5
P1	4	2
P2	9	2


Process	Max - need	Allocation
P0	10	5
P1	4	2
P2	9	<b>3</b>





## Safe, Unsafe, Deadlock State

- If a system is in safe state  $\Rightarrow$  no deadlocks
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock
- Avoidance algorithm:
  - Ensure that a system will never enter an unsafe state



## Avoidance Algorithms

- Single instance of a resource type
  - Use a modified resource-allocation graph
- Multiple instances of a resource type
  - Use the Banker's Algorithm






## Modified Resource-Allocation Graph Scheme

- **Claim edge**  $P_i \text{ ----} \rightarrow R_j$  indicates that process  $P_i$  may request resource  $R_j$
- **Request edge**  $P_i \rightarrow R_j$  indicates that process  $P_i$  requests resource  $R_j$ 
  - Claim edge converts to request edge when a process requests a resource
- **Assignment edge**  $R_j \rightarrow P_i$  indicates that resource  $R_j$  was allocated to process  $P_i$ 
  - Request edge converts to an assignment edge when the resource is allocated to the process
- When a resource is released by a process, assignment edge reconverts to a claim edge
- Resources must be claimed **a priori** in the system




## Resource-Allocation Graph

- Modified Resource-Allocation Graph Example



- Unsafe State In Resource-Allocation Graph





## Resource-Allocation Graph Algorithm

- Suppose that process  $P_i$  requests a resource  $R_j$
- The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph

### Banker's Algorithm

- Multiple instances of resources
- Each process must a priori claim maximum use
- When a process requests a resource, it may have to wait
- When a process gets all its resources it must return them in a finite amount of time



## Data Structures for the Banker's Algorithm

Let  $n$  = number of processes, and  $m$  = number of resources types.

- **Available:** Vector of length  $m$ . If available  $[j] = k$ , there are  $k$  instances of resource type  $R_j$  available
- **Max:**  $n \times m$  matrix. If  $\text{Max}[i,j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$
- **Allocation:**  $n \times m$  matrix. If  $\text{Allocation}[i,j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$
- **Need:**  $n \times m$  matrix. If  $\text{Need}[i,j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task

$$\text{Need}[i,j] = \text{Max}[i,j] - \text{Allocation}[i,j]$$





## Safety Algorithm

1. Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , respectively.  
Initialize:  
 $\text{Work} = \text{Available}$   
 $\text{Finish}[i] = \text{false for } i = 0, 1, \dots, n-1$
2. Find an  $i$  such that both:  
(a)  $\text{Finish}[i] = \text{false}$   
(b)  $\text{Need}_i \leq \text{Work}$   
If no such  $i$  exists, go to step 4
3.  $\text{Work} = \text{Work} + \text{Allocation}_i$   
 $\text{Finish}[i] = \text{true}$   
go to step 2
4. If  $\text{Finish}[i] == \text{true}$  for all  $i$ , then the system is in a **safe state**



## Example of Banker's Algorithm

- 5 processes  $P_0$  through  $P_4$ ;
- 3 resource types: A (10 instances), B (5 instances), and C (7 instances)
- Snapshot at time  $T_0$ :

	Allocation			Max			Available		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	5	3	3	3	2
$P_1$	2	0	0	3	2	2			
$P_2$	3	0	2	9	0	2			
$P_3$	2	1	1	2	2	2			
$P_4$	0	0	2	4	3	3			

Need			
A	B	C	
7	4	3	4
1	2	2	1
6	0	0	5
0	1	1	2
4	3	1	3

	Allocation	Need	Work
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	3 3 2
$P_1$	2 0 0	1 2 2	5 3 2
$P_2$	3 0 2	6 0 0	7 4 3
$P_3$	2 1 1	0 1 1	7 4 5
$P_4$	0 0 2	4 3 1	10 4 7

$$\text{Need}[i,j] = \text{Max}[i,j] - \text{Allocation}[i,j]$$

The system is in a safe state since the sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety criteria





## Expl: based on Need

### Tính toán an toàn dựa vào ma trận Need:

- Khởi tạo  $work=avail \Rightarrow Work(3,3,2)$
- Nếu  $Need \leq Work$ :  $Work = Work + Allo$  (so sánh với Work cập nhật mới nhất theo  $P_i$ )
  - $P0$ :  $Need\_P0(7,4,3) > Work(3,3,2) \Rightarrow$  ko update Work
  - $P1$ :  $Need\_P1(1,2,2) < Work(3,3,2) \Rightarrow Work = Work(3,3,2) + Allo\_P1(2,0,0) \Rightarrow$  update: **Work(5,3,2)**
  - $P2$ :  $Need\_P2(6,0,0) > Work(5,3,2) \Rightarrow$  ko update Work
  - $P3$ :  $Need\_P3(0,1,1) < Work(5,3,2) \Rightarrow Work = Work(5,3,2) + Allo\_P3(2,1,1) \Rightarrow$  update Work(7,4,3)
  - $P4$ :  $Need\_P4(4,3,1) < Work(7,4,3) \Rightarrow Work = Work(7,4,3) + Allo\_P4(0,0,2) \Rightarrow$  update Work(7,4,5)
- Còn  $P0, P2$  chưa cấp phát. Finish = True thì quay lại xem xét cấp phát lại (Step 2)
  - $P0$ :  $Need\_P0(7,4,3) < Work(7,4,5) \Rightarrow Work = Work(7,4,5) + Allo\_P0(0,1,0) \Rightarrow$  update Work(7,5,5)
  - $P2$ :  $Need\_P2(6,0,0) < Work(7,5,5) \Rightarrow Work = Work(7,5,5) + Allo\_P2(3,0,2) \Rightarrow$  update Work(10,5,7)
- Hệ thống trên ở trạng thái an toàn vì nó tồn tại thứ tự an toàn  $P1,3,4,0,2$
- Cách check:
  - So sánh Work ở trạng thái cuối với đề ra có bao nhiêu instant, nếu khớp thì OK
  - Nếu ko cho instant ban đầu thì so Work ở trạng thái cuối với tổng từng kiểu instant ở các Process + instant ở Available, nếu khớp thì OK



## Resource-Request Algorithm for Process $P_i$

**Request<sub>i</sub>** = request vector for process  $P_i$ .

If **Request<sub>i</sub>[j] = k** then process  $P_i$  wants **k** instances of resource type  $R_j$

1. If **Request<sub>i</sub> ≤ Need<sub>i</sub>** go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
2. If **Request<sub>i</sub> ≤ Available**, go to step 3. Otherwise  $P_i$  must wait, since resources are not available
3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:

$$Available = Available - Request_i;$$

$$Allocation_i = Allocation_i + Request_i;$$

$$Need_i = Need_i - Request_i;$$

- If safe  $\Rightarrow$  the resources are allocated to  $P_i$
- If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored

$$Available = Available + Request_i;$$

$$Allocation_i = Allocation_i - Request_i;$$

$$Need_i = Need_i + Request_i;$$





## Example: $P_1$ Request (1,0,2)

- $P_1$  Request (1,0,2): Check that Request  $\leq$  Available (that is,  $(1,0,2) \leq (3,3,2) \Rightarrow$  true

### Before:

	Allocation		
	A	B	C
$P_0$	0	1	0
$P_1$	2	0	0
$P_2$	3	0	2
$P_3$	2	1	1
$P_4$	0	0	2

Available		
A	B	C
3	3	2

Need		
A	B	C
7	4	3
1	2	2
6	0	0
0	1	1
4	3	1

④  
①  
⑤  
②  
③

Avail = Avail - Request;  
Alloc<sub>i</sub> = Alloc<sub>i</sub> + Request;  
Need<sub>i</sub> = Need<sub>i</sub> - Request;

### Update:

	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	4	3	2	3	0
$P_1$	3	0	2	0	2	0			
$P_2$	3	0	2	6	0	0			
$P_3$	2	1	1	0	1	1			
$P_4$	0	0	2	4	3	1			

### Work:

Pro_	A	B	C
cess	2	3	0
$P_1$	5	3	2
$P_3$	7	4	3
$P_4$	7	4	5
$P_0$	7	5	5
$P_2$	10	5	7

- Sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement
  - $\Rightarrow$  R1 can be allocated resource



## Example:

- Can request for (3,3,0) by  $P_4$  be granted?
- Can request for (0,2,0) by  $P_0$  be granted?





## Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme



## Single Instance of Each Resource Type

- a deadlock detection algorithm that uses a variant of the resource-allocation graph, called **wait-for** graph
  - Nodes are processes
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$  to release a resource  $P_i$  needs

We obtain this graph from the resource-allocation graph by removing the resource nodes and collapsing the appropriate edges


- To detect deadlocks:
  - Maintain the wait-for graph
  - Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock
- An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations, where  $n$  is the number of vertices in the graph








## Resource-Allocation Graph and Wait-for Graph



Resource-Allocation Graph



Corresponding wait-for graph



## Several Instances of a Resource Type

- **Available:** A vector of length  $m$  indicates the number of available resources of each type
- **Allocation:** An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process
- **Request:** An  $n \times m$  matrix indicates the current request of each process.
  - If **Request**  $[i][j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .





## Detection Algorithm

1. Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , respectively Initialize:
  - a) **Work** = **Available**
  - b) For  $i = 1, 2, \dots, n$ , if **Allocation<sub>i</sub>**  $\neq 0$ , then **Finish[i] = false**; otherwise, **Finish[i] = true**
2. Find an index  $i$  such that both:
  - a) **Finish[i] == false**
  - b) **Request<sub>i</sub>**  $\leq$  **Work**
 If no such  $i$  exists, go to step 4
3. **Work** = **Work** + **Allocation<sub>i</sub>**  
**Finish[i] = true**  
 go to step 2
4. If **Finish[i] == false**, for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state. Moreover, if **Finish[i] == false**, then  $P_i$  is deadlocked  
 If **Finish[i] == true**, then the system is in safe

*Algorithm requires an order of  $O(m \times n^2)$  operations to detect whether the system is in deadlocked state*



## Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time  $T_0$ :

	Allocation			Request			Available		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	0	0	0	0	0	0
$P_1$	2	0	0	2	0	2			
$P_2$	3	0	3	0	0	0			
$P_3$	2	1	1	1	0	0			
$P_4$	0	0	2	0	0	2			

- Sequence  $\langle P_0, P_2, P_3, P_4, P_1 \rangle$  or
  - Sequence  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$
  - will result in **Finish[i] = true** for all  $i$
- => The system is in safe**

Pro_ cess	A	B	C
P0	0	1	0
P2	3	1	3
P3	5	2	4
P4	5	2	6
P1	7	2	6





## Example (Cont.)

- $P_2$  requests an additional instance of type **C**

Request		
A	B	C
0	0	0
2	0	2
0	0	1
1	0	0
0	0	2

- State of system?
  - Can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes; requests
  - Deadlock exists, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$



## Detection-Algorithm Usage

- When, and how often, to invoke depends on:
  - How often a deadlock is likely to occur?
  - How many processes will need to be rolled back?
    - ▶ One for each disjoint cycle
- If detection algorithm is invoked arbitrarily, there may be many cycles in the resource graph and so we would not be able to tell which of the many deadlocked processes “caused” the deadlock.





## When Deadlock

- When a detection algorithm determines that a deadlock exists, several alternatives are available.
  - One possibility is to inform the operator
  - Another possibility is to let the system recover from the deadlock automatically.
- recover from breaking a deadlock
  - Abort all deadlocked processes
  - Abort one process at a time until the deadlock cycle is eliminated



## Recovery from Deadlock: Process Termination

- In which order should we choose to abort?
  - Priority of the process
  - How long process has computed, and how much longer to completion
  - Resources the process has used
  - Resources process needs to complete
  - How many processes will need to be terminated
  - Is process interactive or batch?





## Recovery from Deadlock: Resource Preemption

- **Selecting a victim** – minimize cost
- **Rollback** – return to some safe state, restart process for that state
- **Starvation** – same process may always be picked as victim, include number of rollback in cost factor



## End of Chapter 2

