



Chapter 4

PROCESS SYNCHRONIZATION

Part 2

Đinh Công Đoàn

1



Content

- Background
- The Critical-Section Problem
- Synchronization Hardware
- Semaphores
- Classical Problems of Synchronization
- Critical Regions
- Monitors

Đinh Công Đoàn

2



Learning Outcomes

- Understand concurrency is an issue in operating systems and multithreaded applications
- Know the concept of a *critical region*.
- Understand how mutual exclusion of critical regions can be used to solve concurrency issues
- Including how mutual exclusion can be implemented correctly and efficiently.
- Be able to identify and solve a *producer consumer bounded buffer* problem.
- Understand and apply standard synchronization primitives to solve synchronisation problems.

Đinh Công Đoàn

3



Background

- Concurrent access to shared data may result in data inconsistency.
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes.
- Shared-memory solution to bounded-buffer problem allows at most $n - 1$ items in buffer at the same time. A solution, where all N buffers are used is not simple.
- Suppose that we modify the producer-consumer code by adding a variable counter, initialized to 0 and incremented each time a new item is added to the buffer

Đinh Công Đoàn

4



Atomic operations

- In concurrent programming, an operation (or set of operations) is **atomic** if it appears to the rest of the system to occur at once without being interrupted.
- Other words used synonymously with atomic are: linearizable, indivisible or uninterruptible.
- Additionally, atomic operations commonly have a succeed-or-fail definition—they either successfully change the state of the system, or have no apparent effect.

Đinh Công Đoàn

5




Non-atomic operations

- When the compiler translates the **shared++** and **shared--** statements, these will be translated to a series of machine instructions (depending on the CPU architecture). On a load/store architecture (for example MIPS):
- **shared++**
 - 1) **load** shared from memory into CPU register
 - 2) **increment** shared and save result in register
 - 3) **store** result back to memory
- **shared--**
 - 1) **load** shared from memory into CPU register
 - 2) **decrement** shared and save result in register
 - 3) **store** result back to memory

Đinh Công Đoàn


6



Interleaving of executing threads

Thread A (increment)				Thread B (decrement)		
OP	Operands	\$t0	BALANCE	OP	Operands	\$t0
			0			
lw	\$t0,	0	0			
addi	\$t0, \$t0, 1	1	0			
sw	\$t0, BALANCE	1	1			
			1	lw	\$t0, BALANCE	1
			1	addi	\$t0, \$t0, -1	0
			0	sw	\$t0, BALANCE	0

If the instructions are executed in this order, the increment and decrement cancel each other and the resulting **BALANCE** is 0.



Interleaving of executing threads

Thread A (increment)				Thread B (decrement)		
OP	Operands	\$t0	BALANCE	OP	Operands	\$t0
			0			
lw	\$t0, BALANCE	0	0			
			0	lw	\$t0, BALANCE	0
addi	\$t0, \$t0, 1	1	0			
			0	addi	\$t0, \$t0, -1	-1
			-1	sw	\$t0, BALANCE	-1
sw	\$t0, BALANCE	1	+1			

Both threads tries to access and update the shared memory location **BALANCE** concurrently. Updates are not atomic and the result depends on the particular order in which the data accesses take place. In this example the resulting **BALANCE** is +1.

Interleaving of executing threads

Thread A (increment)				Thread B (decrement)		
OP	Operands	\$t0	BALANCE	OP	Operands	\$t0
			0			
lw	\$t0, BALANCE	0	0			
			0	lw	\$t0, BALANCE	0
addi	\$t0, \$t0, 1	1	0			
sw	\$t0, BALANCE	1	1			
			1	addi	\$t0, \$t0, -1	-1
			-1	sw	\$t0, BALANCE	-1

Both threads tries to access and update the shared memory location **BALANCE** concurrently. Updates are not atomic and the result depends on the particular order in which the data accesses take place. In this example the resulting **BALANCE** is **-1**.

Concurrency Example

- count is a global variable shared between two threads. After increment and decrement complete, what is the value of count?

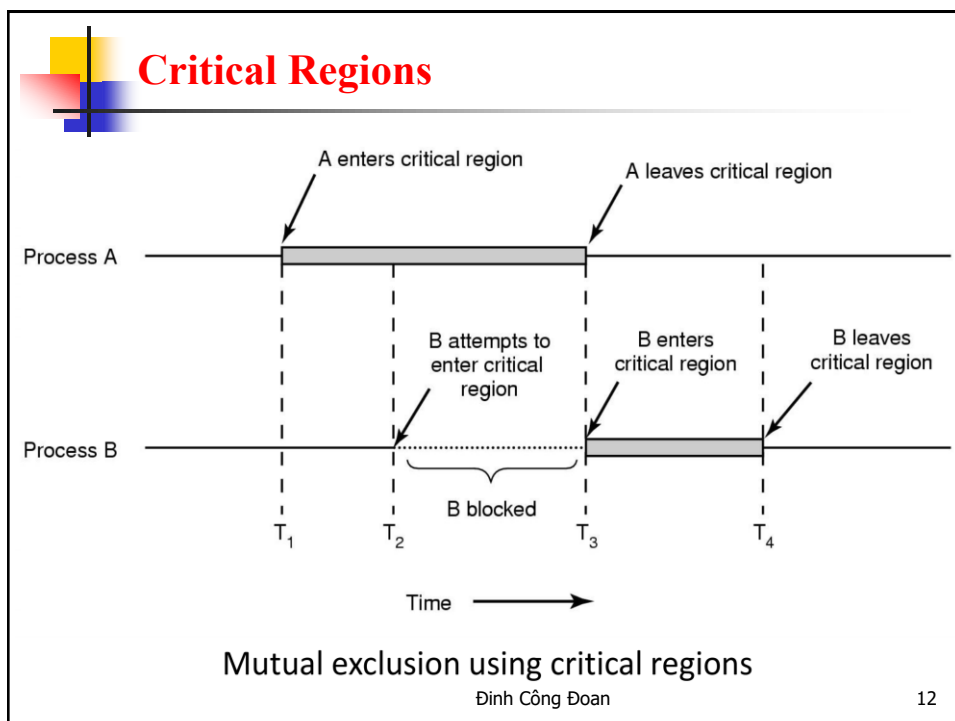
```
void increment ()
{
    int t;
    t = count;
    t = t + 1;
    count = t;
}
```

```
void decrement ()
{
    int t;
    t = count;
    t = t - 1;
    count = t;
}
```

Critical Region

- We can control access to the shared resource by controlling access to the code that accesses the resource. → A *critical region* is a region of code where shared resources are accessed.
 - ✓ Variables, memory, files, etc...
- Uncoordinated entry to the critical region results in a race condition → Incorrect behaviour, deadlock, lost work,...

Đinh Công Đoàn 11





Identifying critical regions

- Critical regions are regions of code that:
 - ✓ Access a shared resource,
 - ✓ and correctness relies on the shared resource not being concurrently modified by another thread/process/entity
- Example critical regions

Đinh Công Đoàn

13



Bounded-Buffer

- Shared data


```
#define BUFFER_SIZE 10
typedef struct {
    ...
} item;
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
int counter = 0;
```

Đinh Công Đoàn

14



Bounded-Buffer

- Producer process


```

item nextProduced;
while (1) {
  while (counter == BUFFER_SIZE)
    ; /* do nothing */
  buffer[in] = nextProduced;
  in = (in + 1) % BUFFER_SIZE;
  counter++;
}
```

Đinh Công Đoàn

15



Bounded-Buffer

- Consumer process


```

item nextConsumed;
while (1) {
  while (counter == 0)
    ; /* do nothing */
  nextConsumed = buffer[out];
  out = (out + 1) % BUFFER_SIZE;
  counter--;
}
```

Đinh Công Đoàn

16



Bounded-Buffer

- The statements
`counter++;`
`counter--;`
 must be performed atomically.
- Atomic operation means an operation that completes in its entirety without interruption
- The statement “**count++**” may be implemented in machine language as:
 - ✓ `register1 = counter`
`register1 = register1 + 1`
`counter = register1`

Đinh Công Đoàn

17



- The statement “**count - -**” may be implemented as:
 - ✓ `register2 = counter`
`register2 = register2 - 1`
`counter = register2`
- If both the producer and consumer attempt to update the buffer concurrently, the assembly language statements may get interleaved.
- Interleaving depends upon how the producer and consumer processes are scheduled.

Đinh Công Đoàn

18



- Assume **counter** is initially 5. One interleaving of statements is
 - ✓ producer: **register1 = counter** (register1 = 5)
 - ✓ producer: **register1 = register1 + 1** (register1 = 6)
 - ✓ consumer: **register2 = counter** (register2 = 5)
 - ✓ consumer: **register2 = register2 - 1** (register2 = 4)
 - ✓ producer: **counter = register1** (counter = 6)
 - ✓ consumer: **counter = register2** (counter = 4)
- The value of **count** may be either 4 or 6, where the correct result should be 5.

Đinh Công Đoàn

19




Race condition

- A race condition or race hazard is the behaviour of an electronic, software or other system where the output is dependent on the sequence or timing of other uncontrollable events.
- It becomes a bug when events do not happen in the intended order.
- The term originates with the idea of two signals racing each other to influence the output first

Đinh Công Đoàn

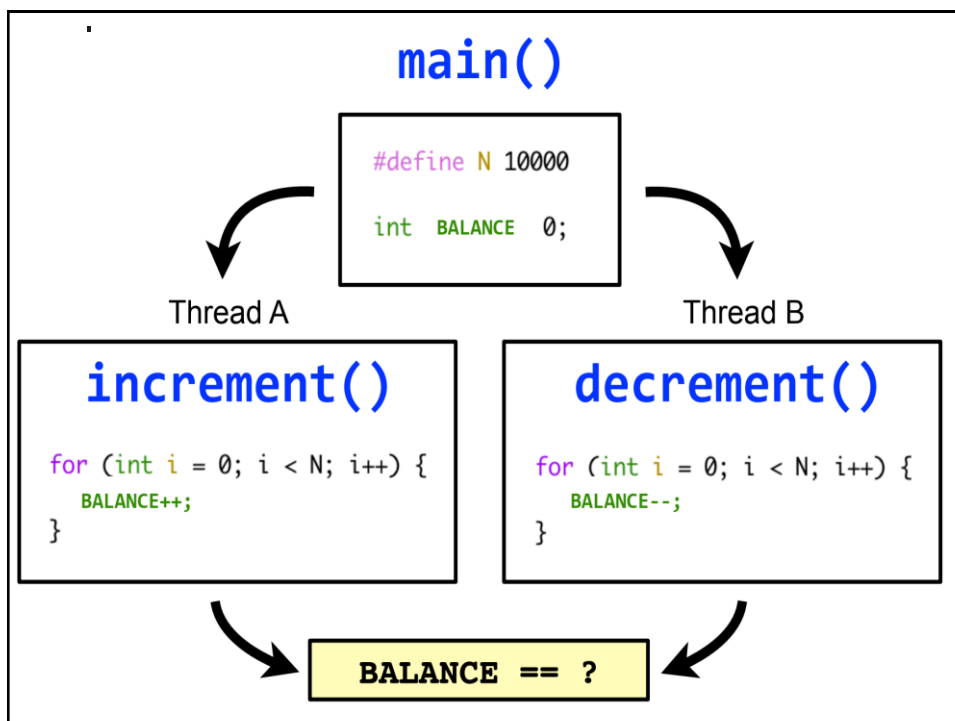
20

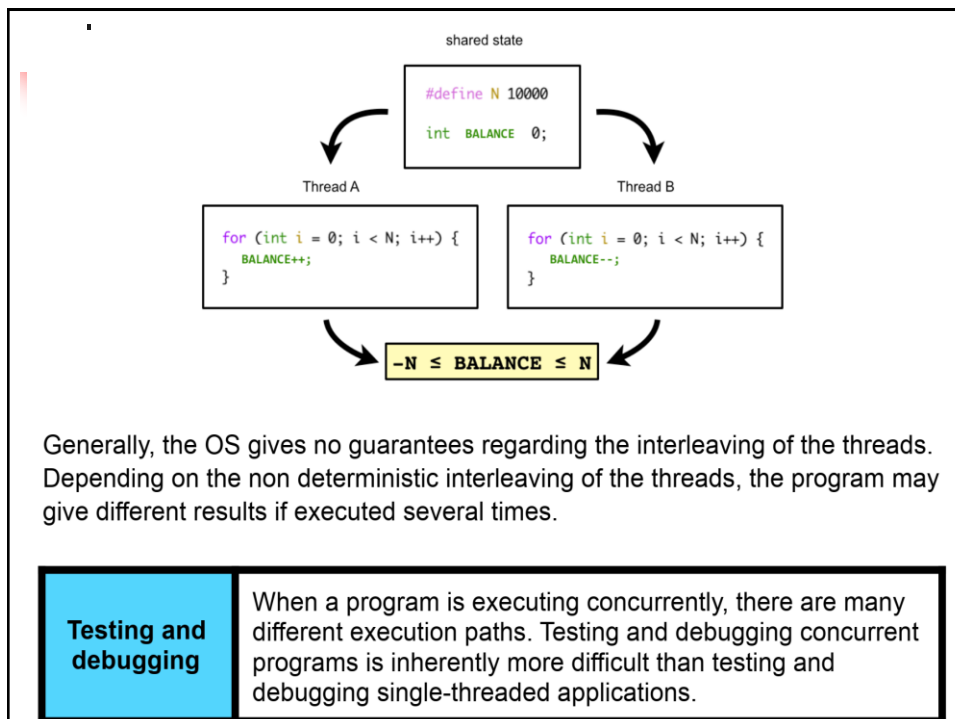
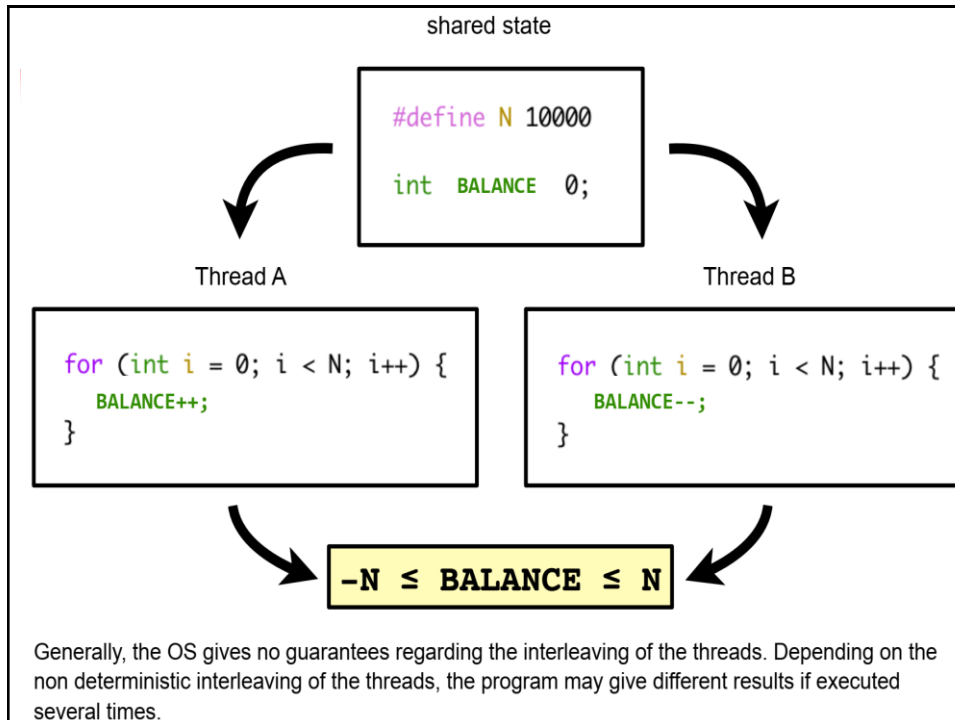



Data race

- A data race occurs when two instructions from different threads access the same memory location and:
 - ✓ at least one of these accesses is a write
 - ✓ and there is no synchronization that is mandating any particular order among these accesses

Đinh Công Đoàn 21








Race condition

- **Race condition:** The situation where several processes access – and manipulate shared data concurrently. The final value of the shared data depends upon which process finishes last.
- To prevent race conditions, concurrent processes must be **synchronized**.
-

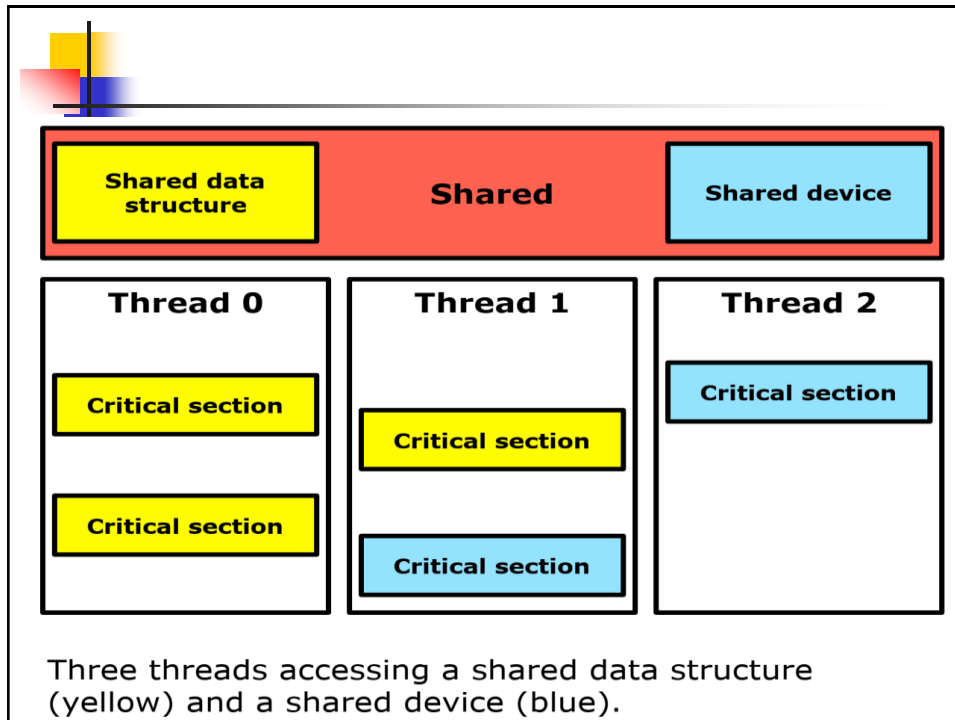
Đinh Công Đoàn 25



Critical section


- Part of a program that should not be concurrently executed by more than one of the program's concurrent processes or threads at a time.
- Typically, the critical section accesses a shared resource, such as a data structure, a peripheral device, or a network connection, that does not allow multiple concurrent accesses.
- By carefully controlling which variables are modified inside and outside the critical section, concurrent access to that state is prevented.

Đinh Công Đoàn 26



The critical – section Problem

- n processes all competing to use some shared data
- Each process has a code segment, called critical section, in which the shared data is accessed.
- Problem – ensure that when one process is executing in its critical section, no other process is allowed to execute in its critical section




- N processes all competing to use shared data.
 - ✓ Structure of process P_i ---- Each process has a code segment, called the critical section, in which the shared data is accessed.

repeat

entry section	/* enter critical section */
critical section	/* access shared variables */
exit section	/* leave critical section */
remainder section	/* do other work */

until false
- Problem
 - ✓ Ensure that when one process is executing in its critical section, no other process is allowed to execute in its critical section

Đinh Công Đoàn 29



Solution to Critical-Section Problem

- 1. **Mutual Exclusion.** If process P_i is executing in its critical section, then no other processes can be executing in their critical sections.
- 2. **Progress.** If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely.
- 3. **Bounded Waiting.** A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted.
 - ✓ Assume that each process executes at a nonzero speed
 - ✓ No assumption concerning relative speed of the n processes.

Đinh Công Đoàn 30

Properties of critical sections

- Assume that each process/thread executes at a nonzero speed. No assumption concerning relative speed of the N processes. In the following discussion the term **task** is used to mean a concurrent unit of execution such as a process or thread

Mutual Exclusion

If task T_i is executing in its critical section, then no other tasks can be executing in their critical sections.

Bounded Waiting

A bound must exist on the number of times that other tasks are allowed to enter their critical sections after a task has made a request to enter its critical section and before that request is granted.

Progress

If no task is executing in its critical section and there exist some task that wish to enter their critical section, then the selection of the task that will enter the critical section next cannot be postponed indefinitely.

More?

Deadlock

Fairness

Starvation

We will come back to this ...

Solution of the Critical Problem

- Software based solutions
- Hardware based solutions
- Operating system based solution
- Structure of Solution

✓ do {


entry section

critical section

exit section

reminder section

✓ } while (1);




Initial attempts to Solve Problem

- Only 2 processes, P0 and P1
- General structure of process P_i (other process P_j)


```

do {
    entry section
    (Primitive) critical section
    exit section
    reminder section
} while (1);
      
```
- Processes may share some common variables to synchronize their actions

Đinh Công Đoàn 33



Algorithm 1 (Dekker 1)

- Shared variables:



```

int turn;
initially turn = 0
turn = i → Pi can enter its critical section
      
```
- Process P_i

```

do {
    while (turn != i) ;
    critical section
    turn = j;
    reminder section
} while (1);
      
```
- Satisfies mutual exclusion, but not progress


Đinh Công Đoàn 34



Algorithm 1

- Satisfies mutual exclusion
 - ✓ The turn is equal to either i or j and hence one of P_i and P_j can enter the critical section
- Does not satisfy progress
 - ✓ Example: P_i finishes the critical section and then gets stuck indefinitely in its remainder section. Then P_j enters the critical section, finishes, and then finishes its remainder section. P_j then tries to enter the critical section again, but it cannot since turn was set to i by P_j in the previous iteration. Since P_i is stuck in the remainder section, turn will be equal to i indefinitely and P_j can't enter although it wants to. Hence no process is in the critical section and hence no progress.
- We don't need to discuss/consider bounded wait when progress is not satisfied

Đinh Công Đoàn 35



Algorithm 2 (Dekker 2)

- Shared variables
boolean flag[2];
 initially **flag [0] = flag [1] = false.**
flag [i] = true → P_i ready to enter its critical section
 Process P_i

```

do {
    flag[i] := true;
    while (flag[j]) ;
    critical section
    flag [i] = false;
    remainder section
} while (1);

```
- Satisfies mutual exclusion, but no progress and no bounded waiting requirement

Đinh Công Đoàn 36



Algorithm 2

- Satisfies mutual exclusion
 - ✓ If P_i enters, then $\text{flag}[i] = \text{true}$, and hence P_j will not enter.
- Does not satisfy progress
 - ✓ Can block indefinitely.... Progress requirement not met
 - ✓ Example: There can be an interleaving of execution in which P_i and P_j both first set their flags to true and then both check the other process' flag. Therefore, both get stuck at the entry section
- We don't need to discuss/consider bounded wait when progress is not satisfied

Đinh Công Đoàn

37



Algorithm 3

- Shared Variables
 - ✓ **var** flag: **array** (0..1) **of** boolean;
initially $\text{flag}[0] = \text{flag}[1] = \text{false}$;
 - ✓ $\text{flag}[i] = \text{true} \Rightarrow P_i$ ready to enter its critical section
- Process P_i

```

repeat
    while  $\text{flag}[j]$  do no-op;
     $\text{flag}[i] := \text{true}$ ;
    critical section
     $\text{flag}[i] := \text{false}$ ;
    remainder section
until false
      
```

Đinh Công Đoàn

38



Algorithm 3

- ***Does not satisfy mutual exclusion***

- ✓ Example: There can be an interleaving of execution in which both first check the other process' flag and see that it is false. Then they both enter the critical section.
- ✓ We don't need to discuss/consider progress and bounded wait when mutual exclusion is not satisfied

Đinh Công Đoàn

39



Algorithm 4 (Peterson)

- Combined shared variables of algorithms 1 and 2.
- Process P_i

```


do {
    flag[i] := true;
    turn = i;
    while (flag[j] and turn = j) ;
    critical section
    flag[i] = false;
    remainder section
} while (1);

```

- Meets all three requirements; solves the critical section problem for two processes.

Đinh Công Đoàn


40



Algorithm 4

- Satisfies mutual exclusion
 - ✓ If one process enters the critical section, it means that either the other process was not ready to enter or it was this process' turn to enter. In either case, the other process will not enter the critical section
- Satisfies progress
 - ✓ If one process exits the critical section, it sets its ready flag to false and hence the other process can enter. Moreover, there is no interleaving in the entry section that can block both.
- Satisfies bounded wait
 - ✓ If a process is waiting in the entry section, it will be able to enter at some point since the other process will either set its ready flag to false or will set to turn to this process.

Đinh Công Đoàn 41



Bakery Algorithm

- Critical section for n processes
 - ✓ Before entering its critical section, process receives a number. Holder of the smallest number enters critical section.
 - ✓ If processes P_i and P_j receive the same number,
 - if $i \leq j$, then P_i is served first; else P_j is served first.
 - ✓ The numbering scheme always generates numbers in increasing order of enumeration; i.e.
1,2,3,3,3,3,4,4,5,5

Đinh Công Đoàn 42

Bakery Algorithm (cont.)

- Notation –

- ✓ Lexicographic order(ticket#, process id#)
 - $(a,b) < (c,d)$ if $(a < c)$ or if $((a = c) \text{ and } (b < d))$
 - $\max(a_0, \dots, a_{n-1})$ is a number, k , such that $k \geq a_i$ for $i = 0, \dots, n-1$

- Shared Data

```
var choosing: array[0..n-1] of Boolean
                    ;(initialized to false)
number: array[0..n-1] of integer
                    ; (initialized to 0)
```

Đinh Công Đoan


43

Bakery Algorithm (cont.)

```

repeat
    choosing[i] := true;
    number[i] := max(number[0], number[1], ..., number[n-1]) + 1;
    choosing[i] := false;
    for j := 0 to n-1
        do begin
            while choosing[j] do no-op;
            while number[j] < 0
                and (number[j], j) < (number[i], i) do no-op;
        end;
    critical section
    number[i] := 0;
    remainder section
until false;

```




Bakery Algorithm

Process	Number
P0	3
P1	0
P2	7
P3	4
P4	8

Đinh Công Đoàn

45




Bakery Algorithm

P0	P2	P3	P4
$(3,0) < (3,0)$	$(3,0) < (7,2)$	$(3,0) < (4,3)$	$(3,0) < (8,4)$
Number[1] = 0	Number[1] = 0	Number[1] = 0	Number[1] = 0
$(7,2) < (3,0)$	$(7,2) < (7,2)$	$(7,2) < (4,3)$	$(7,2) < (8,4)$
$(4,3) < (3,0)$	$(4,3) < (7,2)$	$(4,3) < (4,3)$	$(4,3) < (8,4)$
$(8,4) < (3,0)$	$(8,4) < (7,2)$	$(8,4) < (4,3)$	$(8,4) < (8,4)$

Đinh Công Đoàn


46



Bakery Algorithm

- P1 not interested to get into its critical section □
number[1] is 0
- P2, P3, and P4 wait for P0
- P0 gets into its CS, get out, and sets its number to 0
- P3 get into its CS and P2 and P4 wait for it to get out of its CS
- P2 gets into its CS and P4 waits for it to get out
- P4 gets into its CS
- Sequence of execution of processes: <P0, P3, P2, P4>

Đinh Công Đoàn 47



Bakery Algorithm

- Meets all three requirements:
 - ✓ Mutual Exclusion: $(\text{number}[j], j) < (\text{number}[i], i)$ cannot be true for both P_i and P_j
 - ✓ Bounded-waiting: At most one entry by each process ($n-1$ processes) and then a requesting process enters its critical section (First-Come-First-Serve)
 - ✓ Progress: Decision takes complete execution of the 'for loop' by one process. No process in its 'Remainder Section' (with its number set to 0) participates in the decision making

Đinh Công Đoàn 48

Supporting Synchronization

Programs	Shared Programs
Higher-level API	Locks Semaphores Monitors Send/Receive CRegions
Hardware	Load/Store Disable Ints Test&Set Comp&Swap

- We are going to implement various synchronization primitives using atomic operations
 - ✓ Everything is pretty painful if only atomic primitives are load and store
 - ✓ Need to provide inherent support for synchronization at the hardware level
 - ✓ Need to provide primitives useful at software/user level

Đinh Công Đoàn


49

Synchronization Hardware

- Normally, access to a memory location excludes other accesses to that same location.
- Extension: designers have proposed machine instructions that perform two operations atomically (indivisibly) on the same memory location (e.g., reading and writing).
- The execution of such an instruction is also mutually exclusive (even on Multiprocessors).
- They can be used to provide mutual exclusion but other mechanisms are needed to satisfy the other two requirements of a good solution to the CS problem.

Đinh Công Đoàn

50




Hardware Solutions for Synchronization

- Load/store - Atomic Operations required for synchronization
 - ✓ Showed how to protect a critical section with only atomic load and store ⇒ pretty complex!
- Mutual exclusion solutions presented depend on memory hardware having read/write cycle.
 - ✓ If multiple reads/writes could occur to the same memory location at the same time, this would not work.
 - ✓ Processors with caches but no cache coherency cannot use the solutions
- In general, it is impossible to build mutual exclusion without a primitive that provides some form of mutual exclusion.
 - ✓ How can this be done in the hardware???
 - ✓ How can this be simplified in software???

51

Đinh Công Đoàn




Synchronization Hardware

- Test and modify the content of a word atomically - Test-and-set instruction
- **function** Test-and-Set (**var** target: boolean): boolean;
begin
 Test-and-Set := target;
 target := true;
end;
- Similarly “SWAP” instruction

52

Đinh Công Đoàn



Mutual Exclusion with Test-and-Set

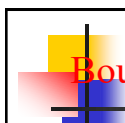
- Shared data: var lock: boolean (initially false)
- Process P_i

```

repeat
    while Test-and-Set (lock) do no-op;
    critical section
    lock := false;
    remainder section
until false;

```


Đinh Công Đoàn 53



Bounded Waiting Mutual Exclusion with Test-and-Set

- **var** $j : 0..n-1$;
 key : boolean;
 repeat
 waiting $[i] := \text{true}$; key := true;
while waiting $[i]$ **and** key **do** key := Test-and-Set(lock);
 waiting $[i] := \text{false}$;
 critical section
 $j := i+1 \bmod n$;
while ($j < i$) **and** (**not** waiting $[j]$) **do** $j := j + 1 \bmod n$;
if $j = i$ **then** lock := false;
else waiting $[j] := \text{false}$;
 remainder section
until false;


Đinh Công Đoàn 54



Solution with TSL

- Is the TSL-based solution good? No
- Mutual Exclusion: Satisfied
- Progress: Satisfied
- Bounded Waiting: Not satisfied

Đinh Công Đoàn 55




Synchronization Hardware

- Atomically swap two variables.

```
void Swap(boolean &a, boolean &b)
{
    boolean temp = a;
    a = b;
    b = temp;
}
```


Đinh Công Đoàn 56



Mutual Exclusion with Swap

- Shared data (initialized to **false**):
 - boolean lock;**
 - boolean waiting[n];**
- 'key' is local and set to false
- Process P_i
 - do {**
 - key = true;**
 - while (key == true)**
 - Swap(lock, key);**
 - critical section
 - lock = false;**
 - remainder section
 - }**


Đinh Công Đoàn 57



Solution with swap

- Is the swap-based solution good? No
- Mutual Exclusion: Satisfied
- Progress: Satisfied
- Bounded Waiting: Not satisfied


Đinh Công Đoàn 58



A Good Solution

- 'key' local; 'lock' and 'waiting' global
- All variables set to false

```
do {
    waiting[i] = true;
    key = true;
    while (waiting[i] && key)
        key = TestAndSet(lock);
    waiting[i] = false;
    Critical Section
}
```




A Good Solution

```

j = (i+1) % n;
while ( (j != i) && !waiting[j] )
    j = (j+1) % n;
if (j == i)
    lock = false;
else
    waiting[j] = false;
    Remainder Section
}

```


Đinh Công Đoàn 60



Solution with Test-And-Set

- Is the given solution good? Yes
- Mutual Exclusion: Satisfied
- Progress: Satisfied
- Bounded Waiting: Satisfied

Đinh Công Đoàn 61



Semaphore


- Semaphore S - integer variable (non-negative)
 - ✓ used to represent number of abstract resources
- Can only be accessed via two indivisible (atomic) operations

wait (S): **while S <= 0 do no-op**
 S := S-1;

signal (S): S := S+1;

 - ✓ P or wait used to acquire a resource, waits for semaphore to become positive, then decrements it by 1
 - ✓ V or signal releases a resource and increments the semaphore by 1, waking up a waiting P, if any
 - ✓ If P is performed on a count <= 0, process must wait for V or the release of a resource.
- **P():“proberen” (to test) ; V() “verhogen” (to increment) in Dutch**


Đinh Công Đoàn 62



Example: Critical Section for n Processes

- Shared variables
 - var** mutex: semaphore
 - initially mutex = 1
- Process P_i
 - repeat**
 - wait(mutex);
 - critical section
 - signal (mutex);
 - remainder section
 - until** false

Đinh Công Đoàn 63




Semaphore as a General Synchronization Tool

- Execute B in P_j only after A execute in P_i
- Use semaphore flag initialized to 0
- Code

P_i	P_j
:	:
:	:
:	:
A	wait(flag)
signal(flag)	B


Đinh Công Đoàn 64



Problem...

- Locks prevent conflicting actions on shared data
 - ✓ Lock before entering critical section and before accessing shared data
 - ✓ Unlock when leaving, after accessing shared data
 - ✓ Wait if locked
- All Synchronization involves waiting
 - ✓ **Busy Waiting**, uses CPU that others could use. This type of semaphore is called a spinlock.
 - Waiting thread may take cycles away from thread holding lock (no one wins!)
 - OK for short times since it prevents a context switch.
 - Priority Inversion: If busy-waiting thread has higher priority than thread holding lock \Rightarrow no progress!
 - ✓ Should sleep if waiting for a long time
- For longer runtimes, need to modify P and V so that processes can block and resume

Đinh Công Đoàn 65




Semaphore Implementation

- Define a semaphore as a record


```

type semaphore = record
    value: integer;
    L: list of processes;
end;
```
- Assume two simple operations
 - ✓ block suspends the process that invokes it.
 - ✓ wakeup(P) resumes the execution of a blocked process P

Đinh Công Đoàn 66




Semaphore Implementation(cont.)

- Semaphore operations are now defined as


```
wait (S): S.value := S.value - 1;
    if S.value < 0
    then begin
        add this process to S.L;
        block;
    end;
signal (S): S.value := S.value + 1;
    if S.value <= 0
    then begin
        remove a process P from S.L;
        wakeup(P);
    end;
```

Đinh Công Đoàn 67



Block/Resume Semaphore Implementation

- If process is blocked, enqueue PCB of process and call scheduler to run a different process.
- Semaphores are executed atomically;
 - ✓ no two processes execute wait and signal at the same time.
 - ✓ Mutex can be used to make sure that two processes do not change count at the same time.
 - If an interrupt occurs while mutex is held, it will result in a long delay.
 - Solution: Turn off interrupts during critical section.

Đinh Công Đoàn 68

Deadlock and Starvation

- Deadlock - two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes.

✓ Let S and Q be semaphores initialized to 1

P₀
wait(S);
wait(Q);
 ⋮
signal(S);
signal(Q);

P₁
wait(Q);
wait(S);
 ⋮
signal(Q);
signal(S);

- Starvation- indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.

Đinh Công Đoàn


69

Two Types of Semaphores

- Counting Semaphore - integer value can range over an unrestricted domain.
- Binary Semaphore - integer value can range only between 0 and 1; simpler to implement.
- Can implement a counting semaphore S as a binary semaphore.

Đinh Công Đoàn


70



Classical Problems of Synchronization

- Bounded Buffer Problem
- Readers and Writers Problem
- Dining-Philosophers Problem

Đinh Công Đoàn 71




Bounded Buffer Problem

- Shared data


```

type item = ....;
var buffer: array[0..n-1] of item;
full, empty, mutex : semaphore;
nextp, nextc :item;
full := 0; empty := n; mutex := 1;
      
```


Đinh Công Đoàn 72



Bounded Buffer Problem

- Producer process - creates filled buffers
 - repeat**
 - ...
 - produce an item in nextp
 - ...
 - wait (empty);
 - wait (mutex);
 - ...
 - add nextp to buffer
 - ...
 - signal (mutex);
 - signal (full);
 - until** false;

Đinh Công Đoàn 73



Bounded Buffer Problem

- Consumer process - Empties filled buffers
 - repeat**
 - wait (full);
 - wait (mutex);
 - ...
 - remove an item from buffer to nextc
 - ...
 - signal (mutex);
 - signal (empty);
 - ...
 - consume the next item in nextc
 - ...
 - until** false;


Đinh Công Đoàn 74

Discussion

- ASymmetry?
 - ✓ Producer does: P(empty), V(full)
 - ✓ Consumer does: P(full), V(empty)
- Is order of P's important?
 - ✓ Yes! Can cause deadlock
- Is order of V's important?
 - ✓ No, except that it might affect scheduling efficiency

Đinh Công Đoàn


75



- Motivation: Consider a shared database
 - ✓ Two classes of users:
 - Readers – never modify database
 - Writers – read and modify database
 - ✓ Is using a single lock on the whole database sufficient?
 - Like to have many readers at the same time
 - Only one writer at a time

Đinh Công Đoàn


76



Readers-Writers Problem

- Shared Data
 - var** mutex, wrt: semaphore (=1);
 - readcount: integer (= 0);
- Writer Process
 - wait(wrt);
 - ...
 - writing is performed
 - ...
 - signal(wrt);


Đinh Công Đoàn 77



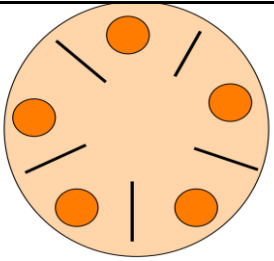
Readers-Writers Problem

- Reader process
 - wait(mutex);
 - readcount := readcount + 1;
 - if** readcount = 1 **then** wait(wrt);
 - signal(mutex);
 - ...
 - reading is performed
 - ...
 - wait(mutex);
 - readcount := readcount - 1;
 - if** readcount = 0 **then** signal(wrt);
 - signal(mutex);

Đinh Công Đoàn 78



Dining-Philosophers Problem




- Shared Data

```

var chopstick: array [0..4] of semaphore (=1
  ■ Philosopher i :
    repeat
      wait (chopstick[i]);
      wait (chopstick[i+1 mod 5]);
      ...
      eat
      ...
      signal (chopstick[i]);
      signal (chopstick[i+1 mod 5]);
      ...
      think
      ...
    until false;
  
```

Đinh Công Đoàn
79



Higher Level Synchronization

- Timing errors are still possible with semaphores
- Example 1


```

signal (mutex);
...
critical region
...
wait (mutex);
      
```
- Example 2



```

wait(mutex);
...
critical region
...
wait (mutex);
      
```
- Example 3


```

wait(mutex);
...
critical region
...
Forgot to signal
      
```


Đinh Công Đoàn
80



Motivation for Other Sync. Constructs

- Semaphores are a huge step up from loads and stores
 - ✓ Problem is that semaphores are dual purpose:
 - They are used for both mutex and scheduling constraints
 - Example: the fact that flipping of P's in bounded buffer gives deadlock is not immediately obvious. How do you prove correctness to someone?
- Idea: allow manipulation of a shared variable only when condition (if any) is met – **conditional critical region**
- Idea : Use locks for mutual exclusion and condition variables for scheduling constraints
- **Monitor:** a lock (for mutual exclusion) and zero or more condition variables (for scheduling constraints) to manage concurrent access to shared data
 - ✓ Some languages like Java provide this natively

Đinh Công Đoàn 81



Conditional Critical Regions

- High-level synchronization construct
- A shared variable *v* of type *T* is declared as:
var v: shared T
- Variable *v* is accessed only inside statement **region v when B do S**
 - ✓ where *B* is a boolean expression.
 - ✓ While statement *S* is being executed, no other process can access variable *v*

Đinh Công Đoàn 82



Critical Regions (cont.)

- Regions referring to the same shared variable exclude each other in time.
- When a process tries to execute the region statement,
 - region** v **when** B **do** S
 - ✓ the Boolean expression B is evaluated.
 - ✓ If B is true, statement S is executed.
 - ✓ If it is false, the process is delayed until B becomes true and no other process is in the region associated with v

Đinh Công Đoan

83



Example - Bounded Buffer

- Shared variables
 - var** buffer: **shared record**
 - pool: **array**[0..n-1] **of** item;
 - count, in, out: integer;
 - end;**
- Producer Process inserts nextp into the shared buffer
 - region** buffer **when** count < n
 - do begin**
 - pool[in] := nextp;
 - in := in+1 **mod** n;
 - count := count + 1;
 - end;**

Đinh Công Đoan

84

Bounded Buffer Example


- Consumer Process removes an item from the shared buffer and puts it in nextc
- **region** buffer **when** count > 0
 - do begin**
 - nextc := pool[out];
 - out := out+1 **mod** n;
 - count := count -1;
 - end;**

Đinh Công Đoàn

85

Monitors

- High-level synchronization construct that allows the safe sharing of an abstract data type among concurrent processes
- **type** monitor-name = **monitor**
 - variable declarations
 - procedure entry** P1 (...);
 - begin ... end;**
 - procedure entry** P2 (...);
 - begin ... end;**
 - ...
 - procedure entry** Pn(...);
 - begin ... end;**
 - begin**
 - initialization code
 - end**




Đinh Công Đoàn

86

Monitor with Condition Variables

- Lock: the lock provides mutual exclusion to shared data
 - ✓ Always acquire before accessing shared data structure
 - ✓ Always release after finishing with shared data
 - ✓ Lock initially free
- Condition Variable: a queue of threads waiting for something inside a critical section
 - ✓ Key idea: make it possible to go to sleep inside critical section by atomically releasing lock at time we go to sleep



Đinh Công Đoàn

87

Monitors with condition variables

- To allow a process to wait within the monitor, a condition variable must be declared, as:

var x,y: condition

 - ✓ Condition variable can only be used within the operations wait and signal. Queue is associated with condition variable.
 - ✓ The operation

x.wait;

 means that the process invoking this operation is suspended until another process invokes

x.signal;
 - ✓ The x.signal operation resumes exactly one suspended process. If no process is suspended, then the signal operation has no effect

Đinh Công Đoàn

88



Dining Philosophers

```

■ type dining-philosophers= monitor
  var state: array[0..4] of (thinking, hungry, eating);
  var self: array[0..4] of condition;
  // condition where philosopher I can delay himself when
  hungry
  but // is unable to obtain chopstick(s)
  procedure entry pickup (i :0..4);
  begin
    state[i] := hungry;
    test(i); //test that your left and right neighbors are not eating
    if state [i] <> eating then self [i].wait;
  end;
  procedure entry putdown (i:0..4);
  begin
    state[i] := thinking;
    test (i + 4 mod 5 ); // signal one neighbor
    test (i + 1 mod 5 ); // signal other neighbor
  end;

```

Đinh Công Đoàn

89



Dining Philosophers (cont.)

```

■ procedure test (k :0..4);
  begin
    if state [k + 4 mod 5] <> eating
      and state [k ] = hungry
      and state [k + 1 mod 5] <> eating
    then
      begin
        state[k] := eating;
        self [k].signal;
      end;
    end;
  begin
    for i := 0 to 4
      do state[i] := thinking;
    end;

```

Đinh Công Đoàn

90

