



Chapter 7

VIRTUAL MEMORY

Đinh Công Đoàn

1




Contents

- Background
- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples

Đinh Công Đoàn


2



Objectives

- To describe the benefits of a virtual memory system
- To explain the concepts of demand paging, page-replacement algorithms, and allocation of page frames
- To discuss the principle of the working-set model

Đinh Công Đoàn 3




Background

- **Virtual memory** – separation of user logical memory from physical memory.
 - ✓ Only part of the program needs to be in memory for execution
 - ✓ Logical address space can therefore be much larger than physical address space
 - ✓ Allows address spaces to be shared by several processes
 - ✓ Allows for more efficient process creation
- Virtual memory can be implemented via:
 - ✓ Demand paging
 - ✓ Demand segmentation


Đinh Công Đoàn 4

Virtual Memory That is Larger Than Physical Memory




Virtual-address Space

Shared Library Using Virtual Memory



Đình Công Đoàn


6



Demand Paging

- Bring a page into memory only when it is needed
 - ✓ Less I/O needed
 - ✓ Less memory needed
 - ✓ Faster response
 - ✓ More users
- Page is needed \Rightarrow reference to it
 - ✓ invalid reference \Rightarrow abort
 - ✓ not-in-memory \Rightarrow bring to memory
- **Lazy swapper** – never swaps a page into memory unless page will be needed
 - ✓ Swapper that deals with pages is a **pager**

Đinh Công Đoàn 7

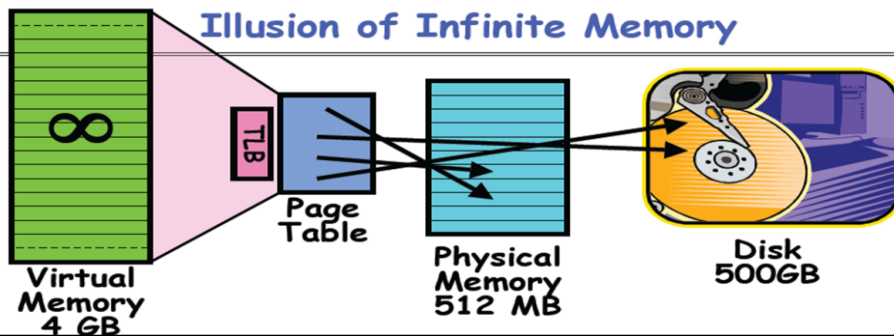


- Modern programs require a lot of physical memory
 - ✓ Memory per system growing faster than 25%-30% per year
- But they don't use all their memory all the time
 - ✓ 90-10 rule: programs spend 90% of their time in 10% of their code
 - ✓ Wasteful to require all of user's code to be in memory
- Solution: use main memory as a cache for disk

Đinh Công Đoàn 8


Demand Paging


- Disk is larger than physical memory, so in-use virtual memory can be bigger than physical memory
 - ✓ Combined memory of running processes can also be much larger than physical memory
 - ✓ More programs fit into memory, allowing more concurrency



Demand Paging

- Principle: transparent level of indirection (page table)
 - ✓ Supports flexible placement of physical data
 - Data could be on disk, or even across network
 - ✓ Variable location of data transparent to user program
 - Performance issue, not correctness issue






Demand Paging

- Demand Paging is Caching:
 - ✓ What is block size? (1 page)
 - ✓ What is the organization?
 - Fully associative: arbitrary virtual → physical mapping
 - ✓ How do we find a page in the cache?
 - Check TLB, then page-table
 - ✓ What is page replacement policy (LRU, random, etc.)?
 - This is today's topic!
 - ✓ What happens on a miss?
 - Page fault: go to disk, and possibly swap a page out to disk to make room
 - ✓ What happens on a write (write-through, write-back)
 - Definitely write-back! Need a dirty bit!

Đinh Công Đoàn 11



Demand Paging

- PTE helps us implement paging
 - ✓ Valid/invalid bit indicates if page is in memory or on disk
- Suppose user references page with invalid PTE?
 - ✓ MMU traps to OS as a page fault
 - ✓ What does the OS do?
 - Choose an old page to replace
 - If old page is dirty, write contents back to disk
 - Invalidate PTE and any cached TLB entry for old page
 - Load new page into memory from disk
 - Update PTE
 - Continue thread from original faulting location
 - ✓ TLB for new page will be loaded when thread continues
 - ✓ While pulling pages off disk for one process, OS runs another process from ready queue
 - After all, the faulting process is essentially blocking on disk I/O
 - Suspended process sits on wait queue

Đinh Công Đoàn 12


Review: Software-Loaded TLB

- MIPS/SIMICS/Nachos TLB is loaded by software
 - ✓ High TLB hit rate: ok to trap to software to fill the TLB, even if slower
 - ✓ Simpler hardware and added flexibility: software can maintain translation tables in whatever convenient format
- How can a process run without access to page table?
 - ✓ Fast path (TLB hit with valid = 1)
 - Translation to physical page done by hardware
 - ✓ Slow path (TLB hit with valid = 0 or TLB miss)
 - Hardware receives a TLB fault
 - ✓ What does the OS do on a TLB Fault?
 - Traverse page table to find appropriate PTE
 - If valid = 1, load PTE into TLB, continue thread
 - If valid = 0, perform Page Fault, then continue thread
- Everything is transparent to the user process
 - ✓ It doesn't know about paging to/from disk
 - ✓ It doesn't even know about software TLB handling

Đinh Công Đoàn

13

Transfer of a Paged Memory to Contiguous Disk Space



Đinh Công Đoàn

14

Valid-Invalid Bit

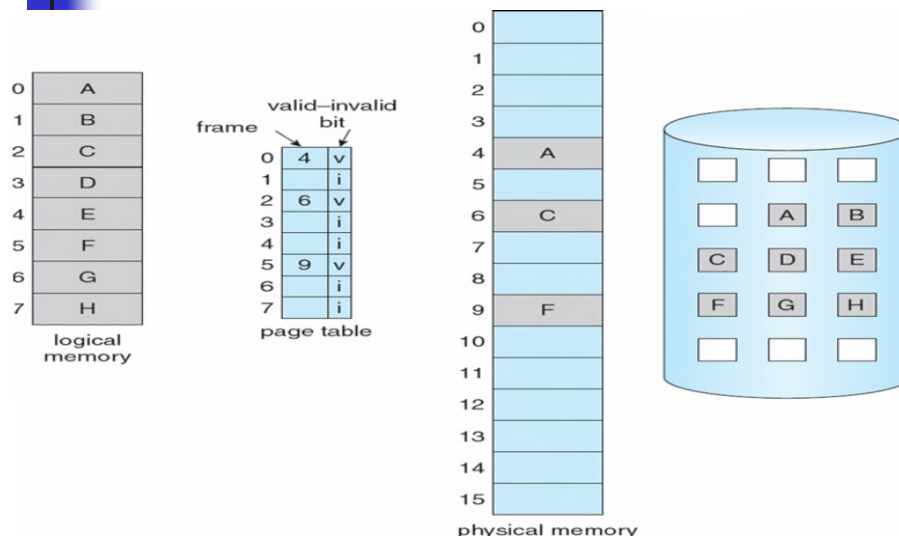
- With each page table entry a valid-invalid bit is associated ($v \Rightarrow$ in-memory, $i \Rightarrow$ not-in-memory)
- Initially valid-invalid bit is set to i on all entries
- Example of a page table snapshot:
- During address translation, if valid-invalid bit in page table entry is $I \Rightarrow$ page fault

	v
	v
	v
	v
	i
....	
	i
	i

Đinh Công Đoàn

15

Page Table When Some Pages Are Not in Main Memory



Đinh Công Đoàn

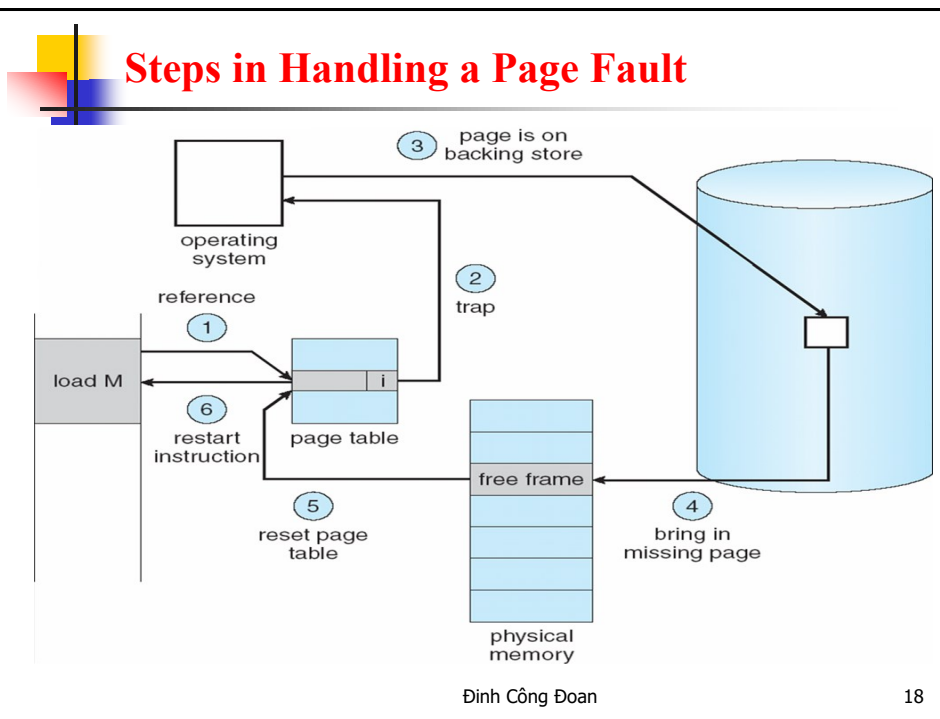
16

Page Fault

- If there is a reference to a page, first reference to that page will trap to operating system: **page fault**
 1. Operating system looks at another table to decide:
 - ✓ Invalid reference \Rightarrow abort
 - ✓ Just not in memory
 2. Get empty frame
 3. Swap page into frame
 4. Reset tables
 5. Set validation bit = **v**
 6. Restart the instruction that caused the page fault

Đinh Công Đoàn

17



Đinh Công Đoàn

18

Re-Run the Faulting Instruction?

- Yes, but what if the instruction has side effects?
 - ✓ Options: Unwind side-effects or finish them off?
 - ✓ Example: `mov (sp)+, 10`
 - What if page fault occurs when writing to the stack pointer? Was `sp` already incremented?
 - ✓ Example: `strcpy (r1), (r2)`
 - Source and destination overlap: can't unwind in principle!
 - IBM S/370 and VAX solution: execute twice – once read-only
 - ✓ Example: RISC delayed branches
 - `div r1, r2, r3`
 - `ld r1, sp`
 - Need to track PC and nPC
 - ✓ Example: Delayed Exceptions
 - `div r1, r2, r3`
 - `ld r1, sp`
 - Takes many cycles to detect divide-by-zero, but `ld` caused the page fault

Đinh Công Đoàn


19

Precise Exceptions

- State of the machine is preserved as if the program executed up to the offending instruction
 - ✓ All previous instructions completed
 - ✓ Offending instruction and all following instructions act as if they have not even started
 - ✓ Same system code will work on different implementations
 - ✓ Difficult in the presence of pipelining, out-of-order execution, etc.
 - ✓ MIPS takes this position
- Imprecise: system software figures out where it is and puts everything back together
- Performance goals often lead designers to forsake precise interrupts (unfortunately)
- Modern techniques for out-of-order execution and branch prediction help implement precise interrupts

Đinh Công Đoàn

20




Performance of Demand Paging

- Page Fault Rate $0 \leq p \leq 1.0$
 - ✓ if $p = 0$ no page faults
 - ✓ if $p = 1$, every reference is a fault
- Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p (\text{page fault overhead} \\ & \quad + \text{swap page out} \\ & \quad + \text{swap page in} \\ & \quad + \text{restart overhead}) \end{aligned}$$

Đinh Công Đoàn 21




Demand Paging Example

- Memory access time = 200 nanoseconds
- Average page-fault service time = 8 milliseconds
- $\text{EAT} = (1 - p) \times 200 + p (8 \text{ milliseconds})$

$$\begin{aligned} &= (1 - p) \times 200 + p \times 8,000,000 \\ &= 200 + p \times 7,999,800 \end{aligned}$$
- If one access out of 1,000 causes a page fault, then
EAT = 8.2 microseconds.
- This is a slowdown by a factor of 40!!


Đinh Công Đoàn 22



Demand Paging

- Does software-loaded TLB need a use bit?
- Two options
 - ✓ Hardware sets use bit in TLB
 - When TLB entry is replaced, software copies use bit back to page table
 - ✓ Software manages TLB entries as FIFO list
 - Everything not in TLB is Second-Chance list, managed as strict LRU
- Core Map
 - ✓ Page tables map virtual page → physical page
 - ✓ Do we need a reverse mapping (i.e. physical page → virtual page)?
 - Yes, clock algorithm runs through page frames. If sharing, then multiple virtual-pages per physical frame can exist
 - Can't push page out to disk without invalidating all of the PTEs

Đinh Công Đoàn 23



Process Creation

- Virtual memory allows other benefits during process creation:
 - ✓ Copy-on-Write
 - ✓ Memory-Mapped Files (later)
- Copy-on-Write
 - ✓ Copy-on-Write (COW) allows both parent and child processes to initially *share* the same pages in memory
 - ✓ If either process modifies a shared page, only then is the page copied
 - ✓ COW allows more efficient process creation as only modified pages are copied
 - ✓ Free pages are allocated from a pool of zeroed-out pages

Đinh Công Đoàn 24



Copy – on - write

- Copy-on-Write (COW) allows both parent and child processes to initially share the same pages in memory.
- If either process modifies a shared page, only then is the page copied.
- COW allows more efficient process creation as only modified pages are copied
- Free pages are allocated from a pool of zeroed-out pages

Đinh Công Đoàn

25

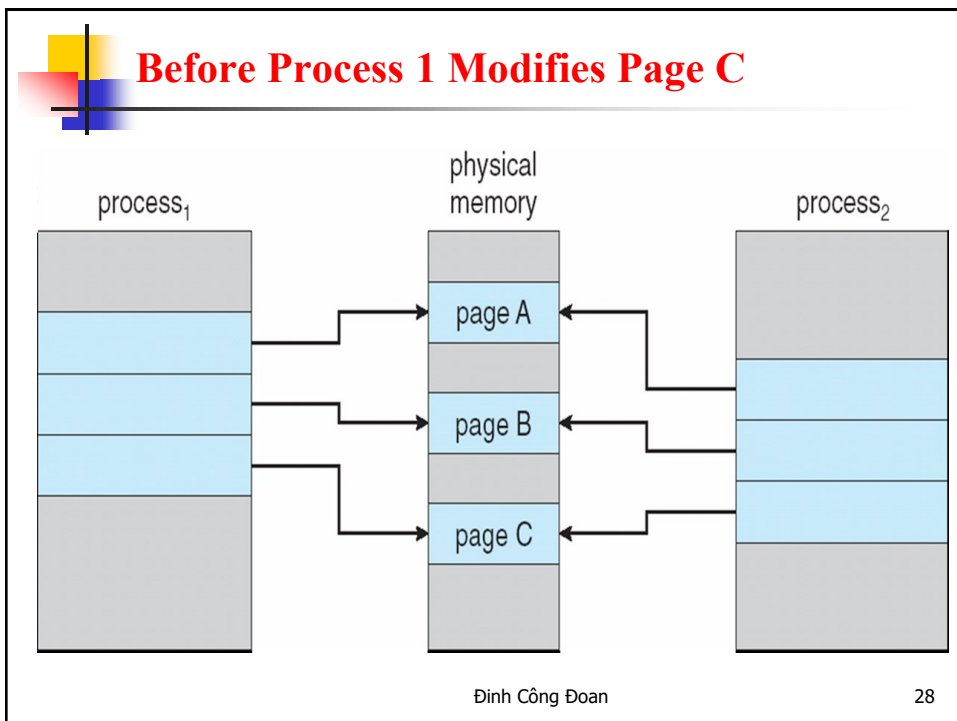
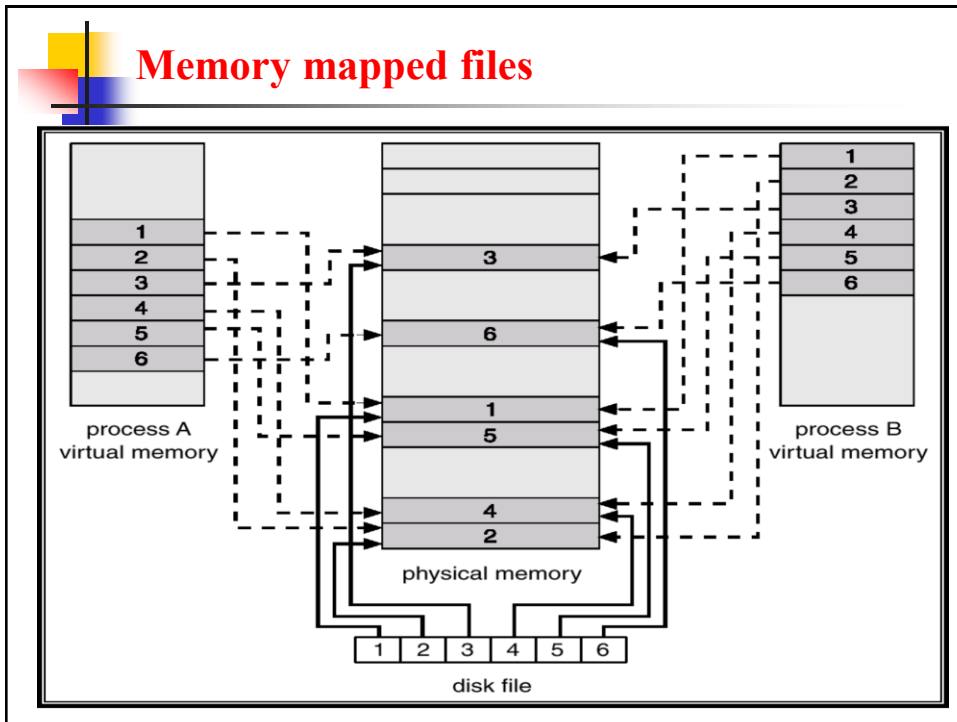


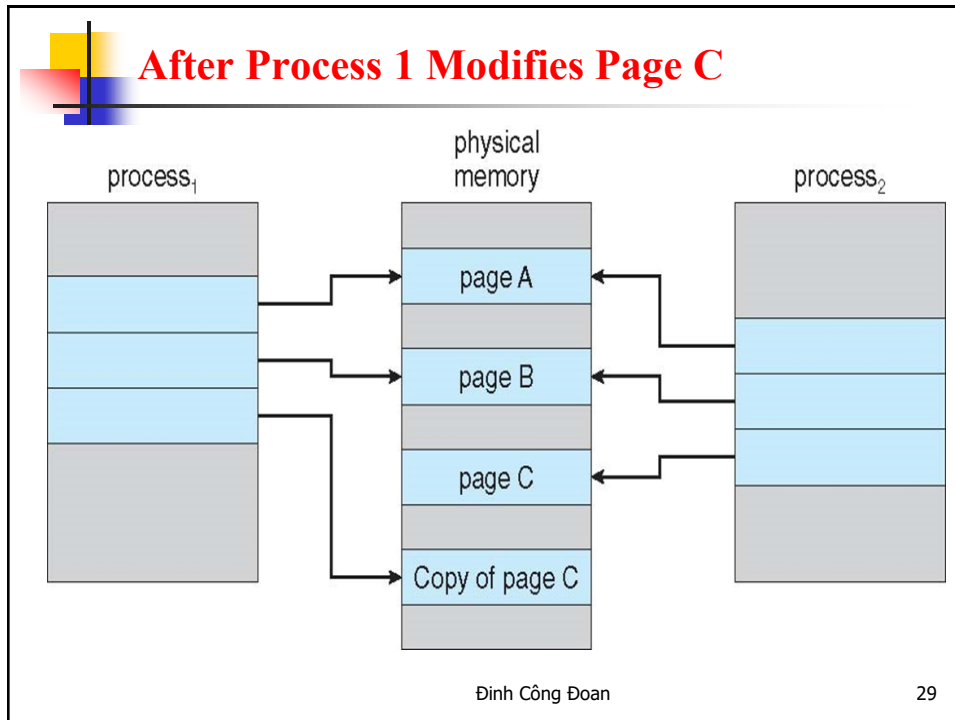
Memory – mapped Files

- Memory-mapped file I/O allows file I/O to be treated as routine memory access by mapping a disk block to a page in memory.
- A file is initially read using demand paging. A pagesized portion of the file is read from the file system into a physical page. Subsequent reads/writes to/from the file are treated as ordinary memory accesses.
- Simplifies file access by treating file I/O through memory rather than **read()** **write()** system calls.
- Also allows several processes to map the same file allowing the pages in memory to be shared

Đinh Công Đoàn

26





What happens if there is no free frame?

- Page replacement – find some page in memory, but not really in use, swap it out
 - ✓ algorithm
 - ✓ performance – want an algorithm which will result in minimum number of page faults
- Same page may be brought into memory several times

Đinh Công Đoàn

30



Page Replacement


- Prevent over-allocation of memory by modifying page-fault service routine to include page replacement
- Use **modify (dirty) bit** to reduce overhead of page transfers – only modified pages are written to disk
- Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory

Đinh Công Đoàn

31



Need For Page Replacement



Đinh Công Đoàn

32



Basic Page Replacement


1. Find the location of the desired page on disk
2. Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim** frame
3. Bring the desired page into the (newly) free frame; update the page and frame tables
4. Restart the process

Đinh Công Đoàn

33




Page Replacement



Đinh Công Đoàn


34



Page Replacement Algorithms

- Replacement is a performance issue: the cost of being wrong is high (going to disk)
- Keep “important” pages in memory, not toss them out
- FIFO
 - ✓ Throw out oldest page. Be fair – let every page live in memory for same amount of time
 - ✓ Bad, because throws out heavily used pages instead of infrequently used ones
- MIN
 - ✓ Replace page that won’t be used for the longest time
 - ✓ Great, but can’t know the future
 - ✓ Used for comparison
- RANDOM
 - ✓ Simple: typical solution for TLB’s
 - ✓ Unpredictable, makes it hard to make real-time guarantees

Đinh Công Đoàn 35



Page Replacement Algorithms


- Want lowest page-fault rate
- Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
- In all our examples, the reference string is

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

Đinh Công Đoàn 36

Graph of Page Faults Versus The Number of Frames

- Ideally, adding memory, the miss rate should go down (above)
- Is this always the case?
 - ✓ It would seem so!
- Belady's Anomaly: Some replacement algorithms don't have this obvious property
 - ✓ FIFO




Đinh Công Đoàn


37

First-In-First-Out (FIFO) Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 frames (3 pages can be in memory at a time per process)



- 4 frames



- Belady's Anomaly: more frames \Rightarrow more page faults

Đinh Công Đoàn

38

FIFO Page Replacement


reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2	2	4	4	4	0		0	0		7	7	7
	0	0	0		3	3	3	2	2	2		1	1		1	0	0
		1	1		1	0	0	0	3	3		3	2		2	2	1

page frames

Đinh Công Đoàn 39



Belady's Anomaly

Ref:	A	B	C	D	A	B	E	A	B	C	D	E
Page:												
1	A			D			E					
2		B			A					C		
3			C			B					D	

Ref:	A	B	C	D	A	B	E	A	B	C	D	E
Page:												
1	A						E				D	
2		B						A				E
3			C						B			
4				D						C		

Đinh Công Đoàn 41

MIN: Optimal Algorithm


- Replace page that will not be used for longest period of time
- 4 frames example

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1
2
3
4

- How do you know this?
- Used for measuring how well your algorithm performs

Đinh Công Đoàn 42



MIN: Optimal Page Replacement


reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2			2			2						7		
	0	0	0		0		4		0		0						0		
		1	1		3		3		3		1						1		

page frames

Đinh Công Đoàn 43




Least Recently Used (LRU) Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1	1	1	1	5
2	2	2	2	2
3	5	5	4	4
4	4	3	3	3

- List Implementation
 - ✓ On each use, remove page from list and place it at head
 - ✓ LRU page is at the tail
- But need to traverse the list to move a page from the middle to the head

Đinh Công Đoàn 44




Least Recently Used (LRU) Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1	1	1	1	5
2	2	2	2	2
3	5	5	4	4
4	4	3	3	3

- Counter implementation
 - ✓ Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
 - ✓ When a page needs to be changed, look at the counters to determine which are to change

Đinh Công Đoàn 45



LRU Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2					2					4	4	4	0					1	1	1
	0	0	0					0					0	0	3	3					3	0	0
		1	1					3					3	2	2	2					2	2	7

page frames

- Stack implementation – keep a stack of page numbers in a double link form:
 - ✓ Page referenced:
 - move it to the top
 - requires 6 pointers to be changed
 - ✓ No search for replacement.

Đinh Công Đoàn 46

■ Use Of A Stack to Record The Most Recent Page References

reference string

4 7 0 7 1 0 1 2 1 2 7 1 2

2
1
0
7
4

stack
before
a

7
2
1
0
4

stack
after
b

↑ a ↑ b

Đinh Công Đoàn

47

Comparing the Algorithms

- FIFO: Suppose we have 3 frames, 4 pages, and the following reference stream

✓ A B C A B D A D B C B

- 7 Faults

✓ When referencing D, replacing A is a bad choice, since we need A again right away

Ref:	A	B	C	A	B	D	A	D	B	C	B
Page:											
1	A					D				C	
2		B					A				
3			C						B		

Đinh Công Đoàn

48

Comparing the Algorithms

- MIN: Suppose we have 3 frames, 4 pages, and the following reference stream
 - ✓ A B C A B D A D B C B
- 5 Faults
 - ✓ D is brought in to the page that will not be referenced for the longest time
- How does this compare to LRU?

Ref:	A	B	C	A	B	D	A	D	B	C	B
Page:											
1	A									C	
2		B									
3			C			D					

Đinh Công Đoàn

49


When Will LRU Perform Poorly?

- Consider
 - ✓ A B C D A B C D A B C D
- Every LRU reference is a page fault!

Ref:	A	B	C	D	A	B	C	D	A	B	C	D
Page:												
1	A			D			C			B		
2		B			A			D			C	
3			C			B			A			D

- Consider
 - ✓ A B C D A B C D A B C D
- MIN does much better


Ref:	A	B	C	D	A	B	C	D	A	B	C	D
Page:												
1	A									B		
2		B					C					
3			C	D								



Least Recently Used (LRU) Algorithm

- Due to locality, it would seem that LRU would be a good candidate to approximate MIN
- But LRU is hard to implement, takes up a lot of space or time.
- In practice, we **approximate** LRU instead.

Đinh Công Đoàn 51




LRU Approximation Algorithms

- Reference bit
 - ✓ With each page associate a bit, initially = 0
 - ✓ When page is referenced bit set to 1
 - ✓ Replace the one which is 0 (if one exists)
 - We do not know the order, however
- Second chance
 - ✓ Need reference bit
 - ✓ Clock replacement
 - ✓ If page to be replaced (in clock order) has reference bit = 1 then:
 - set reference bit 0
 - leave page in memory
 - replace next page (in clock order), subject to same rules

Đinh Công Đoàn 52

Clock Algorithm


- Arrange physical pages in circle with single clock hand
 - ✓ Replace an old page, not the oldest page
 - ✓ Approximation to LRU, hence an approximation to an approximation of MIN
- Details
 - ✓ Hardware “use” bit per physical page
 - Hardware sets use bit on each reference
 - If use bit isn’t set, means not referenced for a long time
 - Nachos hardware sets use bit in the TLB, you have to copy this back to page table when TLB entry gets replaced
 - ✓ On Page Fault
 - Advance clock hand (not real time)
 - Check use bit: 1 == used recently; clear and leave alone, 0 == selected candidate for replacement
 - Even if all use bits set, will eventually loop around; in which case, resort to FIFO



Đinh Công Đoàn

53

Clock Algorithm

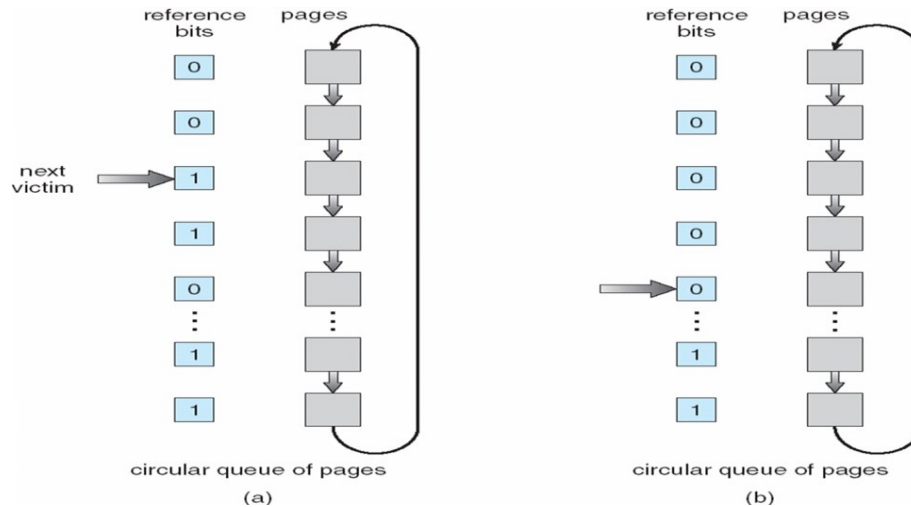


- What if hand moving slowly?
 - ✓ Good sign or bad sign?
 - Not many page faults and/or find page quickly
- What if hand moving quickly?
 - ✓ Lots of page faults and/or lots of reference bits set
- One way to view clock algorithm
 - ✓ Crude partitioning of pages into two groups: young and old
 - ✓ Why not partition into more groups?
 - Clock Algorithm a.k.a. “Second Chance” – could have Nth chance, too!
 - Counting algorithms

Đinh Công Đoàn

54

Second-Chance (clock) Page-Replacement Algorithm



Đinh Công Đoàn


55

Counting Algorithms: Nth Chance Version of Clock Algorithm

- Nth Chance Algorithm: OS keeps counter per page:
sweeps
- On page fault, OS checks use bit:
 - ✓ 1 == clear use and also clear counter (used in last sweep)
 - ✓ 0 == increment counter, if count = N, replace page
- So the clock hand has to sweep by N times without page being used before page is replaced
- How do we pick N?
 - ✓ Why pick large N?
 - ✓ Why pick small N?
- What about dirty pages?

Đinh Công Đoàn


56

- 
- Nth Chance Algorithm: OS keeps counter per page: # sweeps
 - On page fault, OS checks use bit:
 - ✓ 1 == clear use and also clear counter (used in last sweep)
 - ✓ 0 == increment counter, if count = N, replace page
 - So the clock hand has to sweep by N times without page being used before page is replaced
 - How do we pick N?
 - ✓ Why pick large N?
 - Better approximation to LRU: If $N \sim 1K$, really good approximation
 - ✓ Why pick small N?
 - Might have to look a long way to find a free page with large N
 - What about dirty pages?

Đinh Công Đoàn


57

Counting Algorithms: Nth Chance Version of Clock Algorithm

- 
- What about Dirty Pages?
 - ✓ Takes extra overhead to replace a dirty page, so give dirty pages an extra chance before replacing?
 - ✓ Common approach
 - Clean pages, use $N = 1$
 - Dirty pages, use $N = 2$, and write back to disk when $N = 1$
 - Counting Algorithms
 - ✓ Keep a counter of the number of references that have been made to each page
 - ✓ **LFU Algorithm**: replaces page with smallest count
 - ✓ **MFU Algorithm**: based on the argument that the page with the smallest count was probably just brought in and has yet to be used

Đinh Công Đoàn


58



Clock Algorithms: Details

- Which bits of a PTE entry are useful to us?
 - ✓ Use: set when a page is referenced; cleared by clock algorithm
 - ✓ Modified (Dirty): set when a page is modified, cleared when a page is written to disk
 - ✓ Valid: ok for program to reference this page
 - ✓ Read-Only: ok for program to read page, but not modify
 - For example, for catching modifications to code pages
- Do we really need hardware-supported modified bit?
 - ✓ No, we can emulate it using a read-only bit

Đinh Công Đoàn 59



Clock Algorithms: Details

- Which bits of a PTE entry are useful to us?
 - ✓ Use: set when a page is referenced; cleared by clock algorithm
 - ✓ Modified (Dirty): set when a page is modified, cleared when a page is written to disk
 - ✓ Valid: ok for program to reference this page
 - ✓ Read-Only: ok for program to read page, but not modify
 - For example, for catching modifications to code pages
- Do we really need hardware-supported modified bit?
 - ✓ No, we can emulate it (BSD Unix) using a read-only bit
 - Initially, mark all pages as read only (even the data pages)
 - On write, trap to the OS, OS software sets software modified bit, and marks page as read-write
 - When page comes back in from disk, mark read-only

Đinh Công Đoàn 60

Clock Algorithms: Details

- Do we really need a hardware-supported use bit?
 - ✓ No, we can emulate it similar to the modified bit
 - Mark all pages as invalid, even if in memory
 - On read to invalid page, trap to OS
 - OS sets use bit, and marks page read-only
 - ✓ Get modified bit in same way as previous
 - On write, trap to OS (either invalid or read-only)
 - Set use and modified bits, mark page read-write
 - ✓ When clock hand passes by, reset use and modified bits and mark page as invalid again
- Remember, however, that clock is just an approximation of LRU
 - ✓ Can we do a better approximation, given that we have to take page faults on some reads and writes to collect use information?
 - ✓ Need to identify old page, not the oldest page!
 - ✓ Answer: second chance list

Đinh Công Đoàn

61

Second-Chance List Algorithm

- VAX/VMS approach
- Split memory in two: Active List (RW), Second Chance list (Invalid)
- Access pages in Active List at full speed
- Otherwise, Page Fault
 - ✓ Move overflow page from end of Active List to front of Second Chance list
 - ✓ Mark overflow page invalid
 - ✓ If desired page is on the Second Chance list:
 - Move to front of Active List
 - Mark RW
 - ✓ Else:
 - Page in to front of Active List
 - Mark RW
 - Page out LRU victim at the end of the Second Chance list

Đinh Công Đoàn

62



Second-Chance List Algorithm

- How many pages for second chance list?
 - ✓ If 0 → FIFO
 - ✓ If all → LRU, but page fault on every reference
- Pick intermediate value. Result is:
 - ✓ Few disk accesses (page only goes to disk if unused for a long time) (+)
 - ✓ Increased overhead trapping to OS (software/hardware tradeoff) (-)
- With page translation, we can adapt to any kind of access the program makes
 - ✓ Can use page translation / protection to share memory between threads on widely separated machines
- Why didn't VAX just include a "use" bit?
 - ✓ Strecker (architect) asked OS people, they said they didn't need it, so didn't implement it
 - ✓ He later got blamed, but VAX did ok anyway

Đinh Công Đoàn

63




Allocation of Frames

- Each process needs *minimum* number of pages
- Example: IBM 370 – 6 pages to handle SS MOVE instruction:
 - ✓ instruction is 6 bytes, might span 2 pages
 - ✓ 2 pages to handle *from*
 - ✓ 2 pages to handle *to*
- Two major allocation schemes
 - ✓ fixed allocation
 - ✓ priority allocation

Đinh Công Đoàn

64

Free List



- Keep set of free frames ready for use in demand paging
 - ✓ Free list filled in background by clock algorithm or other technique (clock daemon)
 - ✓ Dirty pages start copying back to disk when they enter the list
- Like VAX second-chance list
 - ✓ If page needed before reused, just return to the active set
- Advantage: faster for page fault
 - ✓ Can always use page (or pages) immediately on fault

Đinh Công Đoàn

65


Fixed Allocation

- Equal allocation – For example, if there are 100 frames and 5 processes, give each process 20 frames.
- Proportional allocation – Allocate according to the size of process

 $s_i = \text{size of process } p_i$
 $S = \sum s_i$
 $m = \text{total number of frames}$
 $a_i = \text{allocation for } p_i = \frac{s_i}{S} \times m$
 $m = 64$
 $s_i = 10$
 $s_2 = 127$
 $a_1 = \frac{10}{137} \times 64 \approx 5$
 $a_2 = \frac{127}{137} \times 64 \approx 59$

Đinh Công Đoàn


66



Priority Allocation

- Use a proportional allocation scheme using priorities rather than size
- If process P_i generates a page fault,
 - ✓ select for replacement one of its frames
 - ✓ select for replacement a frame from a process with lower priority number

Đinh Công Đoàn 67




Global vs. Local Replacement

- **Global replacement** – process selects a replacement frame from the set of all frames; one process can take a frame from another
- **Local replacement** – each process selects from only its own set of allocated frames
- Priority Allocation may, if using global replacement, select for replacement a frame from a process with lower priority.

Đinh Công Đoàn 68

Adaptive Allocation based on Page Fault Frequency

- Can we reduce capacity misses by dynamically changing the number of pages per application?



- Establish “acceptable” page fault rate
 - ✓ If rate is too low, process loses a frame to a process that needs it
 - ✓ If rate is too high, process gains a frame
- But what if we just don’t have enough memory?

Đinh Công Đoàn


69

Thrashing

- If a process does not have “enough” pages, the page-fault rate is very high. This leads to:
 - ✓ low CPU utilization
 - ✓ operating system thinks that it needs to increase the “degree of multiprogramming” (number of ready processes)
 - ✓ another process added to the ready queue
 - Which requires more memory, leading to more page faults, leading to lower CPU utilization, and the cycle continues...
- **Thrashing** \equiv a process is busy swapping pages in and out
- How do we detect and respond to thrashing?

Đinh Công Đoàn

70



Demand Paging and Thrashing

- Why does demand paging work? Locality model
 - ✓ Process migrates from one locality to another
 - ✓ Localities may overlap
- Why does thrashing occur? Σ size of locality $>$ total memory size

Đinh Công Đoàn

72



Locality In A Memory-Reference Pattern

- Program memory access patterns have temporal and spatial locality
 - ✓ Group of pages accessed along a given time slice is called the “Working Set”
 - ✓ Working Set defines the minimum number of pages needed for the process to “behave well”
- Not enough memory for the working set → Thrashing
 - ✓ Better to swap out process?

Đinh Công Đoàn

73



Working-Set Model

page reference table


... 2 6 1 5 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...



- $\Delta \equiv$ working-set window \equiv a fixed number of page references
Example: 10,000 instruction
- WSS_i (working set of Process P_i) = total number of pages referenced in the most recent Δ (varies in time)
 - ✓ if Δ too small will not encompass entire locality
 - ✓ if Δ too large will encompass several localities
 - ✓ if $\Delta = \infty \Rightarrow$ will encompass entire program
- $D = \sum WSS_i \equiv$ total demand frames
- if $D > m \Rightarrow$ Thrashing
- Policy if $D > m$, then suspend one of the processes

Đinh Công Đoàn


74



What About Compulsory Misses?

- Recall that compulsory misses are misses that occur the first time that a page is seen
 - ✓ Page that are touched for the first time
 - ✓ Pages that are touched after process is swapped out/swapped back in
- Clustering:
 - ✓ On a page fault, bring in multiple pages “around” the faulting page
 - ✓ Since efficiency of disk reads increases with sequential reads, makes sense to read several sequential pages
- Working Set Tracking
 - ✓ Use algorithm to try to track working set of an application
 - ✓ When swapping process back in, swap in working set


Đinh Công Đoàn 75



Keeping Track of the Working Set

- Approximate with interval timer + a reference bit
- Example: $\Delta = 10,000$
 - ✓ Timer interrupts after every 5000 time units
 - ✓ Keep in memory 2 bits for each page
 - ✓ Whenever a timer interrupts copy and sets the values of all reference bits to 0
 - ✓ If one of the bits in memory = 1 \Rightarrow page in working set
- Why is this not completely accurate?
- Improvement = 10 bits and interrupt every 1000 time units


Đinh Công Đoàn 76



Paging Summary

- Replacement Policies
 - ✓ FIFO
 - ✓ MIN
 - ✓ LRU
 - Clock Algorithm
 - Nth Chance Clock Algorithm
 - Second-Chance List Algorithm
- Working Set
 - ✓ Set of pages touched by a process recently
- Thrashing
 - ✓ A process is busy swapping pages in and out
 - ✓ Process will thrash if working set doesn't fit in memory
 - ✓ Need to swap out a process

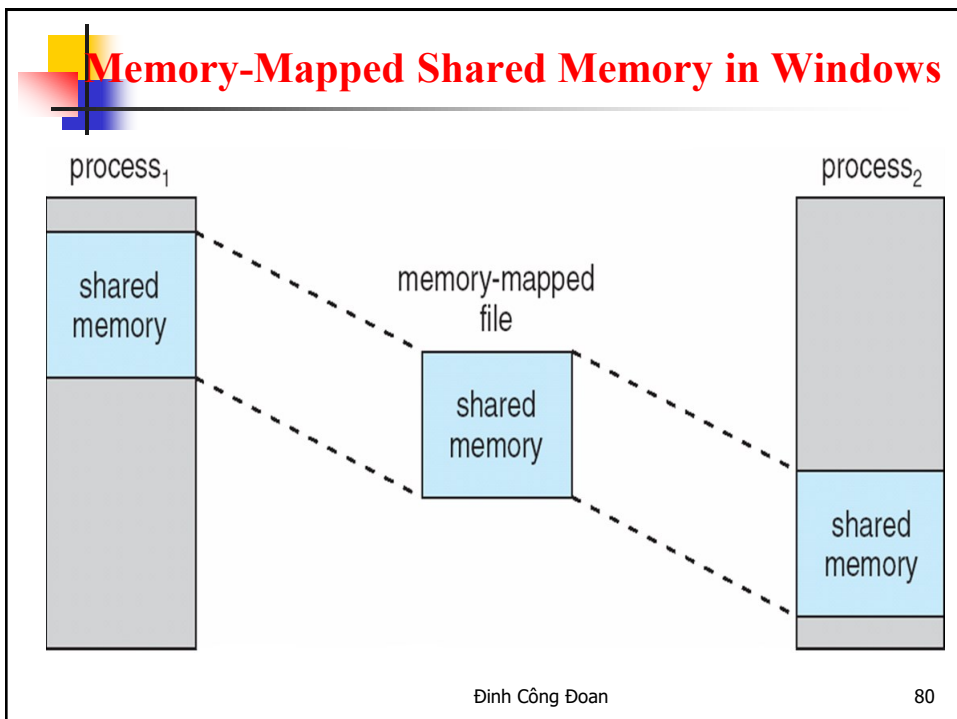
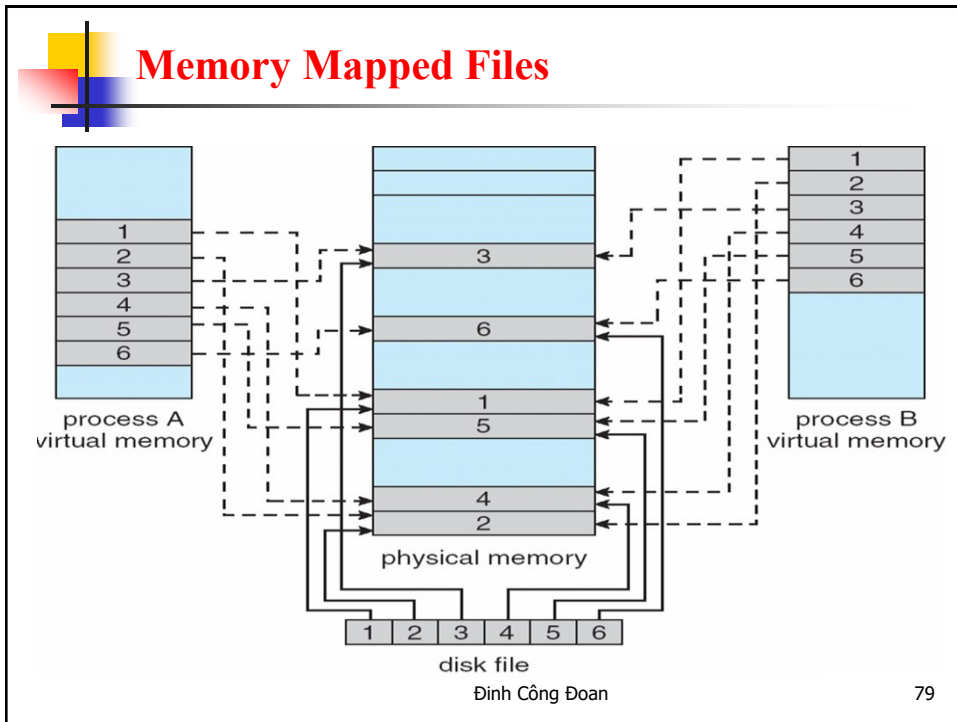
Đinh Công Đoàn 77




Memory-Mapped Files

- Memory-mapped file I/O allows file I/O to be treated as routine memory access by **mapping** a disk block to a page in memory
- A file is initially read using demand paging. A page-sized portion of the file is read from the file system into a physical page. Subsequent reads/writes to/from the file are treated as ordinary memory accesses.
- Simplifies file access by treating file I/O through memory rather than `read()` `write()` system calls
- Also allows several processes to map the same file allowing the pages in memory to be shared

Đinh Công Đoàn 78






Allocating Kernel Memory

- Treated differently from user memory
- Often allocated from a free-memory pool
 - ✓ Kernel requests memory for structures of varying sizes
 - ✓ Some kernel memory needs to be contiguous

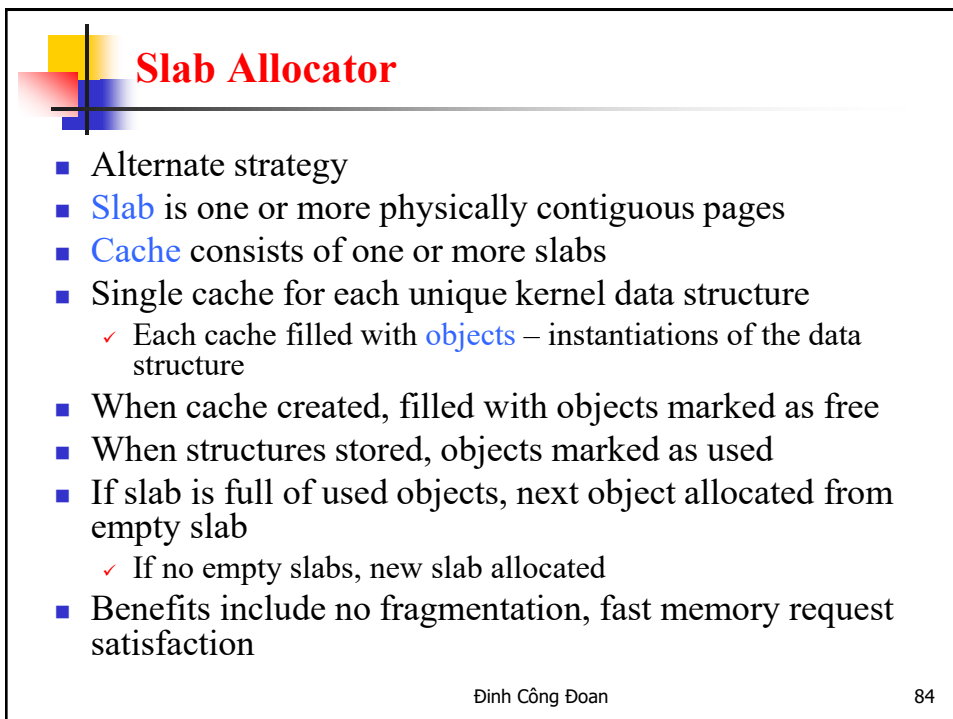
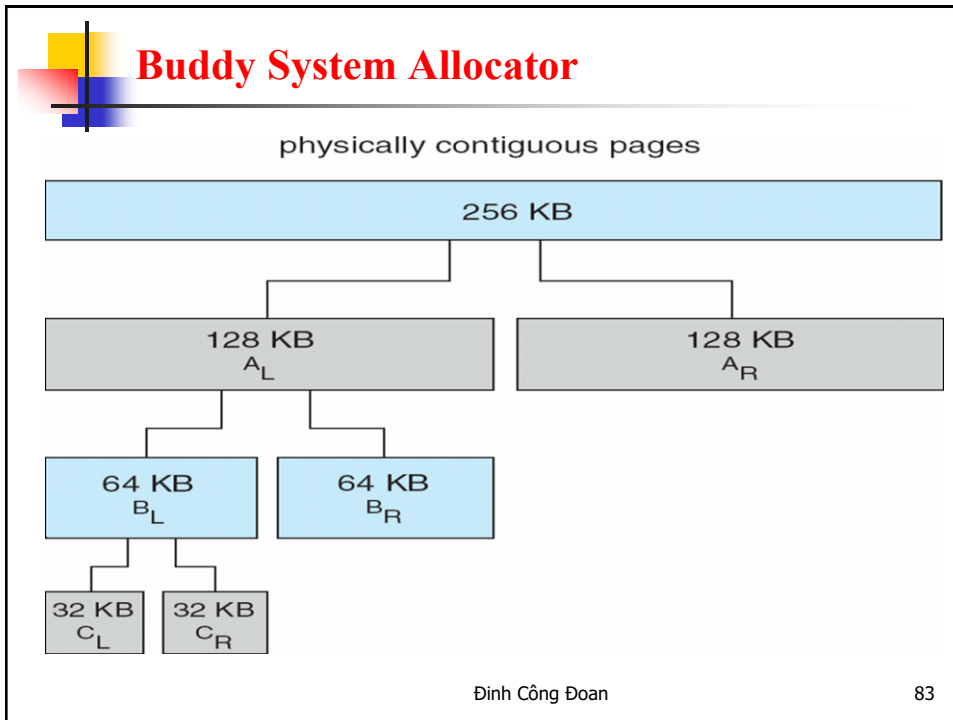
Đinh Công Đoàn 81

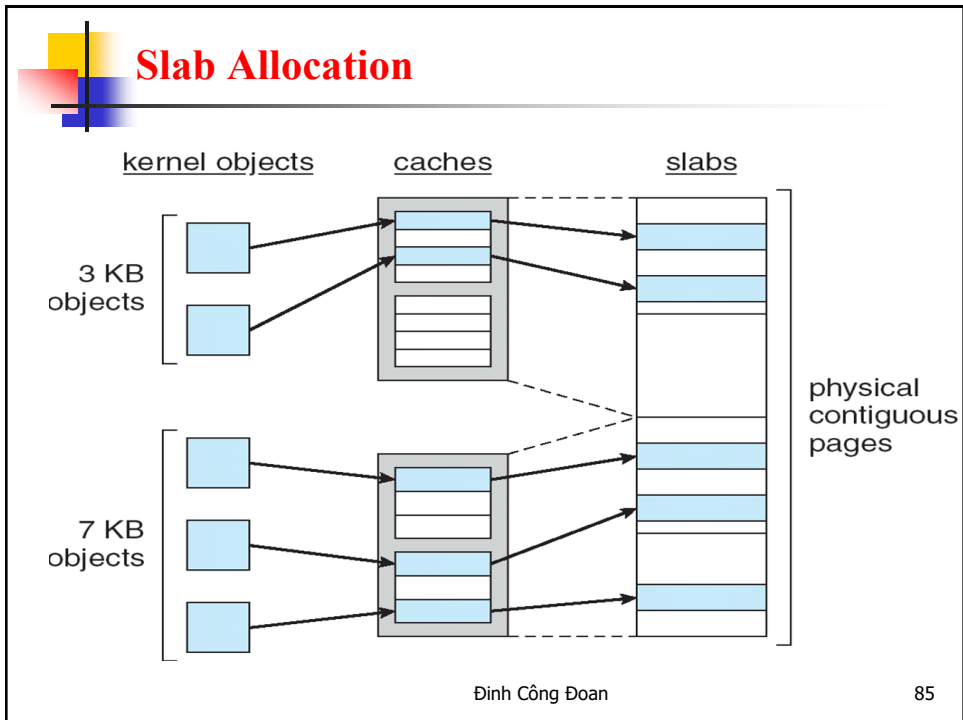


Buddy System

- Allocates memory from fixed-size segment consisting of physically-contiguous pages
- Memory allocated using power-of-2 allocator
 - ✓ Satisfies requests in units sized as power of 2
 - ✓ Request rounded up to next highest power of 2
 - ✓ When smaller allocation needed than is available, current chunk split into two buddies of next-lower power of 2
 - Continue until appropriate sized chunk available

Đinh Công Đoàn 82





Other Issues -- Prepaging

- **Prepaging**
 - ✓ To reduce the large number of page faults that occurs at process startup
 - ✓ Prepage all or some of the pages a process will need, before they are referenced
 - ✓ But if prepaged pages are unused, I/O and memory was wasted
 - ✓ Assume s pages are prepaged and α of the pages is used
 - Is cost of $s * \alpha$ save pages faults > or < than the cost of prepaging
 - $s * (1 - \alpha)$ unnecessary pages?
 - α near zero \Rightarrow prepaging loses

Đinh Công Đoàn

86



Other Issues – Page Size

- Page size selection must take into consideration:
 - ✓ fragmentation
 - ✓ table size
 - ✓ I/O overhead
 - ✓ locality

Đinh Công Đoàn

87



Other Issues – TLB Reach

- TLB Reach - The amount of memory accessible from the TLB
- $\text{TLB Reach} = (\text{TLB Size}) \times (\text{Page Size})$
- Ideally, the working set of each process is stored in the TLB
 - ✓ Otherwise there is a high degree of page faults
- Increase the Page Size
 - ✓ This may lead to an increase in fragmentation as not all applications require a large page size
- Provide Multiple Page Sizes
 - ✓ This allows applications that require larger page sizes the opportunity to use them without an increase in fragmentation

Đinh Công Đoàn

88

Other Issues – Program Structure

- Program structure

- ✓ `Int[128,128] data;`
- ✓ Each row is stored in one page
- ✓ Program 1


```
for (j = 0; j < 128; j++)
  for (i = 0; i < 128; i++)
    data[i,j] = 0;
```

$128 \times 128 = 16,384$ page faults

- ✓ Program 2

```
for (i = 0; i < 128; i++)
  for (j = 0; j < 128; j++)
    data[i,j] = 0;
```

128 page faults

Đinh Công Đoàn

89


Other Issues – I/O interlock

- **I/O Interlock** – Pages must sometimes be locked into memory
- Consider I/O - Pages that are used for copying a file from a device must be locked from being selected for eviction by a page replacement algorithm

Đinh Công Đoàn

90

Reason Why Frames Used For I/O Must Be In Memory



Đinh Công Đoàn

91